



Universal Knowledge Graph (UKG) MVP – System Documentation

System Overview

The **Universal Knowledge Graph (UKG) MVP** is an AI-driven knowledge management and reasoning system that unifies structured knowledge with advanced simulation capabilities. Its primary purpose is to consolidate complex, multi-domain information into a **universal knowledge graph** and enable intelligent “what-if” reasoning on that data ¹. In essence, UKG acts as a digital twin of an organization’s knowledge: it aggregates data from disparate sources into a connected graph and allows AI agents to run simulations and answer complex queries using that graph ¹. This empowers users to explore scenarios and get decision support in a sandbox-like environment, leveraging the entirety of the knowledge base.

High-Level Architecture: The UKG MVP architecture centers around three core components – a **13-axis coordinate engine**, an **AI simulation stack**, and a **unified knowledge database (USKD)**:

- **13-Dimensional Coordinate System:** All knowledge in UKG is indexed in a **13-axis coordinate space**, meaning every knowledge item (e.g. a fact, document clause, or data point) is assigned a unique 13-dimensional coordinate ². Each axis corresponds to a different context (such as knowledge domain, industry sector, time, location, etc.), so a coordinate situates an item across multiple taxonomies simultaneously ³ ⁴. This coordinate engine is the backbone that **maps any query or data into the shared knowledge space**. (For example, a regulatory clause could have a coordinate encoding its domain, industry code, regulatory context, relevant roles, location, and time – all in one identifier ⁴.)
- **Simulation Stack:** The reasoning engine is implemented as a **multi-layer AI simulation stack** that processes queries through a sequence of cognitive steps. It uses a fixed **10-layer reasoning pipeline** – from initial question parsing to final answer validation – to break down and simulate the thought process in stages ⁵. Within this stack, a **Quad Persona** module runs parallel reasoning from four expert perspectives (more below), and iterative refinement loops improve the answer until high confidence is achieved. This simulation engine allows UKG to not just retrieve facts, but to **simulate multi-perspective reasoning and scenario analysis** on the knowledge graph ⁶.
- **Universal Simulated Knowledge Database (USKD):** The USKD is the **knowledge graph repository** that stores all the content (nodes) and relationships in memory for fast access during simulation ⁷. It acts as the unified knowledge base of the system, aggregating data from various sources into a single graph structure ⁸. The MVP leverages an in-memory graph for performance, while also maintaining a persistent storage layer for durability. Together, UKG and USKD form an end-to-end platform: **the USKD provides the knowledge substrate, and the simulation engine operates on that substrate to answer queries** ⁸.

Overall, the UKG MVP is designed to ensure that any complex question can be mapped onto the knowledge graph and answered by **AI “thinking” through the relevant dimensions** of context. Executives might use it to get a 360° view of a scenario, while engineers can rely on its robust, deterministic reasoning for high-assurance answers. This document provides a comprehensive breakdown of the system’s architecture and modules so engineering teams can implement the MVP from scratch.

Technical Architecture

Backend Architecture (Python FastAPI)

The backend of UKG is implemented in **Python**, built around a FastAPI web service that exposes the system’s functionality via RESTful (and optionally, MPC) APIs. FastAPI was chosen for its asynchronous performance and intuitive integration with Python’s data science stack. The backend is organized into several layers:

- **API Layer:** A FastAPI server defines endpoints for core operations (e.g. querying the knowledge graph, managing knowledge data, configuring personas). Incoming requests (such as a query from a user) are first handled by this layer. The API layer performs **authentication** and basic request validation (e.g. checking query payload format), then hands off to the reasoning core. FastAPI’s asynchronous features allow the server to handle multiple concurrent queries efficiently, and if needed, long-running tasks can be handed to background workers (e.g. via Celery or an async task queue) so the API can respond immediately with a job ID.
- **Service Layer (Coordinators):** This layer contains the orchestrators for major functionalities:
 - The **Query Coordinator** parses and routes queries into the 13-axis engine (invoking the Query Engine module described later). It orchestrates the multi-step simulation: calling the coordinate mapping, triggering knowledge retrieval, running the simulation stack, and collecting the results.
 - The **Knowledge Base Service** interfaces with the USKD data layer for any retrieval or update operations on the knowledge graph. For example, when the Query Engine requests all nodes relevant to “Healthcare regulations in 2025”, this service fetches those from the graph (possibly via a query to a graph database or an in-memory lookup).
 - The **Persona/Role Manager** manages the Quad Persona contexts. It instantiates persona objects (one per role) when a simulation starts and helps coordinate their “memories” (context data) and results.
 - The **Configuration Manager** handles system settings (axes definitions, thresholds, toggles). This would load any custom axis definitions or simulation parameters (like recursion depth limits, confidence thresholds) and provide them to other components.
- **Core Logic Layer:** The core algorithms of UKG reside here, implemented as Python classes or modules:
 - The **13-Axis Coordinate Engine** module provides utilities for interpreting and constructing coordinates. It contains the formal schema for the 13 axes (for example, definitions of Pillar codes, industry code mappings, etc.), and functions to encode or decode a coordinate string. It also has

logic for hierarchical coordinate comparisons (to tell if one coordinate is a sub-node of another) using the Nuremberg notation (dot notation).

- The **Simulation Engine** module implements the 10-layer reasoning process and recursive refinement workflow. This is effectively the “brain” of UKG. It may be structured as a state machine or pipeline where each layer is a function or class (Layer1 through Layer10) that transforms the state. The Simulation Engine uses multiple **Knowledge Algorithms (KAs)** – discrete algorithmic components for specific reasoning tasks (search, inference, validation, etc.) ⁹. These KAs are invoked at appropriate layers. For example, a “*Query Parsing KA*” might be used in Layer1 to interpret the question, while a “*Consistency Check KA*” runs in Layer10 to verify the final answer.
- The **Quad Persona Engine** (or Point-of-View Engine) is a sub-component that the Simulation Engine calls when multi-perspective reasoning is needed (starting at Layer3). It instantiates four **Persona agents** (as separate threads or async tasks in Python) and manages their parallel execution and subsequent aggregation of results.
- The **MCP Integration Module** (if using Model Context Protocol) handles translating internal state to the MCP interface. For instance, if an external tool calls `simulate_scenario` via MCP, this module takes that call, prepares the UKG context (mapping the provided input into axes), triggers the simulation, and then formats the result in MCP’s standard output format.

Given Python’s rich ecosystem, the backend leverages libraries like **NetworkX** for in-memory graph operations and possibly **Pydantic** for data models. Persistence (Neo4j, etc.) is abstracted behind repository classes so the core logic can treat the knowledge graph uniformly. The design emphasizes in-memory operation for speed – the entire reasoning pipeline operates without blocking on external I/O after the initial data load, enabling sub-second responses for many queries ¹⁰ ¹¹.

Example: A `POST /api/query` request handler in FastAPI might accept JSON like `{"question": "...", "user": "...", "params": {...}}`. It authenticates the user, calls `QueryCoordinator.process_query(...)` which uses the coordinate engine to map the question to an initial coordinate vector, queries the USKD for relevant knowledge nodes, runs the Simulation Engine’s main loop, and returns the answer with supporting evidence. Pseudocode for the main flow might look like:

```
@app.post("/api/query")
async def query_ukg(query: QueryRequest):
    user = auth.verify_user(query.token)
    coord = coordinate_engine.map_query_to_coordinate(query.text)
    knowledge_context = knowledge_service.retrieve_context(coord)
    result = simulation_engine.run(knowledge_context, coord)
    return {"answer": result.text, "confidence": result.confidence, "trace": result.explanation}
```

Frontend Architecture (TypeScript and Theming)

The frontend is a web application built with **TypeScript**, using a modern framework (e.g. React) for a rich interactive UI. It adheres to an **enhanced design system** that ensures a consistent look-and-feel across the app. Key aspects of the frontend architecture include:

- **Component-Based Structure:** The UI is composed of reusable components following an *atomic design* hierarchy (atoms, molecules, organisms). For example, basic UI atoms like buttons and inputs are implemented once and used throughout ¹² ¹³. Molecules (combinations of atoms, like a search bar with input and button) and higher-level composites (cards, dialogs) build up the interface. This modular structure (often reflected in a `/components` directory with subfolders for atoms, molecules, etc.) makes the UI extensible and maintainable ¹² ¹⁴.
- **Theming and Design Tokens:** The application supports both **light and dark modes**, using a theming system driven by design tokens. Core style attributes like colors, font sizes, and spacing are centralized (e.g. in a SCSS or CSS-in-JS theme file) so that switching themes is straightforward. The **color palette** defines a primary brand color (deep blue #003366) and secondary accent (teal #008080) ¹⁵. These, along with neutrals and semantic colors, are applied consistently. Typography is standardized (e.g. headers use Roboto Bold, body text Roboto Regular 16px) ¹⁶. Spacing and sizing use a consistent scale (e.g. an 8px baseline grid for margins/padding). The **dark mode** primarily swaps backgrounds and text colors (e.g. dark background #1a1a1a for sidebar ¹⁷, light text), while maintaining the same token values. Theming is often implemented via CSS variables or a context provider that supplies the theme to all components.
- **Layout:** The app uses a responsive layout with a **fixed top navigation bar** (header) and a **collapsible sidebar** for main navigation. The top nav includes branding (logo), user menu, language selector, etc., and is fixed to top ¹⁸. The sidebar (drawer) contains the navigation menu (sections of the app) and can collapse to icons for more screen space ¹⁹. Breadcrumbs are shown below the header to indicate the user's location in the knowledge hierarchy or application section ²⁰. The content area uses a fluid container that adapts to screen size. A **12-column grid system** (based on Bootstrap 5 or similar) is used for arranging content cards and panels, with defined breakpoints for mobile, tablet, desktop responsiveness ²¹. This ensures the UI works on various devices, down to a tablet or phone (the sidebar may turn into a hamburger menu on small screens, etc.).
- **UI Components:** The interface includes a library of custom components to view and interact with the knowledge graph and simulation results:
- **Knowledge Cards:** These are used to display key information snippets (like a summary of a knowledge node or an answer). They have a white background, subtle drop shadow, and rounded corners for a clean card-style look ²². For example, a "knowledge card" might show a regulation clause with its title and a brief description, or a simulation result with the answer and confidence.
- **Sidebar Menu and Search:** The sidebar menu items are styled with appropriate spacing and typography (Roboto 14px) ²³. The sidebar also includes a search input for quickly finding nodes or content, with advanced filtering options (e.g. filter by Pillar or Sector).
- **Dialogs/Modals:** Standard modal dialogs are used for confirmations, settings, or detailed views. They follow the theming (e.g. dark overlay, content box with border-radius and the same shadow

style as cards). Focus management and keyboard accessibility are ensured for modals (so users can tab through and close via keyboard).

- **Forms and Inputs:** Input components (text fields, dropdowns, toggles) are styled consistently (e.g. a 2px border with slight rounding, which turns teal on focus to match the brand). They include proper labels and ARIA attributes for accessibility ²⁴. Form layouts use the grid and spacing system to align labels and fields.
- **Dynamic Elements:** There is an embedded **chat/assistant widget** (perhaps in a corner) for interactive Q&A. This might be implemented as a collapsible chat bubble that expands into a chat window where a user can converse with the AI. It includes an input box and a send button with hover effects (e.g. the send button darkens on hover) ²⁵ ²⁶. Typing indicators and response streaming can be shown to improve UX.
- **Data Visualizations:** If the MVP includes any analytics or graph visualization (e.g. a network graph of related knowledge nodes), it uses libraries like D3.js or Three.js, embedded in the front-end. These visualizations adopt the theme (custom colors, fonts) to stay consistent ²⁷.
- **Accessibility:** The UI is built to **WCAG 2.1 AA** standards. All interactive controls have descriptive text or ARIA labels, ensuring screen readers can describe them ²⁸. The app supports keyboard navigation (e.g. pressing Tab will move through links/buttons in logical order, menus can be opened via keyboard). There is a visible focus indicator for focused elements. Users can toggle a high-contrast mode for better visibility ²⁹. Additionally, the front-end provides options to adjust font size or switch themes for accessibility (e.g. an “accessibility” menu in the header with a font size slider and a dark mode toggle ³⁰).

The front-end communicates with the backend via JSON over HTTPS. It handles showing loading states (using skeleton placeholders on cards, etc. ³¹) while waiting for API responses. Error states are also gracefully handled (e.g. if a query fails, an error message is shown with an option to retry). In summary, the frontend provides an intuitive, modern interface for users to query the UKG and explore results, all built on a solid design system for consistency and maintainability.

System Orchestration and Deployment

The UKG MVP is designed to run in cloud environments using containerization for easy scaling and management. **Docker** containers package the backend (FastAPI server and associated services) and optionally the front-end (if served separately). The typical deployment model is **Kubernetes (K8s)** orchestration, where the UKG backend runs as one or more pods behind a load balancer.

- **Containerization:** All components are containerized for consistency across environments. For instance, a `ukg-backend` container image includes the Python FastAPI app, necessary model files, and any initialization to load the knowledge base. A `ukg-frontend` image might serve the static built front-end (or we use a cloud bucket for static files). These containers can be deployed in a cluster, allowing horizontal scaling (multiple instances) to handle higher query throughput. The documentation includes guidance on Docker and CI/CD setup ³² – for example, building images on each release, and using CI pipelines to run tests and then push images to a registry.
- **Kubernetes & Microservices:** In a Kubernetes deployment, the FastAPI app could run as a deployment with an auto-scaler (scaling based on CPU or request load). If the architecture is split

into microservices (not mandatory for MVP, but possible), we might have separate services for “Knowledge Graph DB”, “Vector Index Service”, and “Simulation Service” communicating via internal APIs or a message bus. However, for the MVP, it’s feasible to keep it as a single service process that includes all logic (monolithic container). K8s will handle service discovery and scaling. A typical K8s setup includes a Service for the API (exposed via an Ingress with TLS), and possibly sidecar containers for logging or monitoring.

- **Stateful Components:** The persistent knowledge graph (if using Neo4j or similar) might be run as its own pod (statefulset or managed DB service) that the UKG backend connects to. Alternatively, if an in-memory only approach is used for MVP, the system might load a snapshot of the knowledge data on startup. In that case, persistence can be achieved by storing snapshots (e.g. JSON or database dumps of knowledge) that the container image or entrypoint loads. For scaling in-memory usage, techniques like sharding by Pillar or using an external distributed cache could be considered later. The MVP likely keeps things simpler: one instance holds the USKD in-memory (assuming the data volume is manageable).
- **Message Queue (if applicable):** If certain operations are long-running (e.g. ingesting a large new dataset into the knowledge graph, or performing an extensive simulation), an asynchronous job queue (like RabbitMQ with Celery, or Kafka) can be introduced. The FastAPI endpoint would enqueue a job and immediately return a job ID, and a worker pod would process the task and store results to a DB or cache for retrieval. This decoupling ensures the main API remains responsive. For typical query simulations (which are optimized to be fast), this may not be needed in MVP, but the architecture allows plugging this in for heavy tasks.
- **Edge and Offline Deployment:** In addition to cloud, UKG MVP can be deployed in an **offline or edge environment** for users who need on-premises or offline capability. In an edge deployment, the system is packaged with a **pre-loaded knowledge graph dataset** and possibly a smaller local model, allowing it to run without internet connectivity. For example, an edge package might include a lightweight container with the FastAPI backend and a database file of the knowledge nodes, so that a field office or a secure facility can run UKG on local machines (with a GPU for acceleration if needed). The edge version would disable any cloud-only features and rely on local compute; queries run entirely on the local hardware (possibly with CUDA support if using large neural network components). The MVP could use SQLite or embedded Neo4j for local storage and embed the necessary model files so that everything runs self-contained. This ensures critical knowledge is accessible even in air-gapped or low-connectivity scenarios.
- **Hybrid Cloud-Edge Model:** In a hybrid deployment, edge instances of UKG run locally but periodically **sync with the cloud**. For example, the cloud-hosted UKG could publish knowledge updates (new nodes or changes in regulations) to a secure endpoint, and edge instances fetch and merge these updates during off-peak hours. The hybrid model might let heavy computations (like training new models or large graph analytics) run in the cloud, while the edge handles real-time querying with latest synced data. A message-passing system or MQTT could handle these syncs. From an orchestration perspective, the edge nodes register with the cloud and the cloud can push critical updates or receive aggregated insights from edge analytics.

All deployments prioritize **secure communications and container security**. Docker images are built with minimal base images, and Kubernetes Pod Security Policies (or equivalent) ensure pods have limited

permissions. Role-based access control on the cluster restricts who can deploy or view data. The deployment guide for the MVP covers setting up the cluster, environment variables for secrets (DB passwords, API keys), and how to monitor the system's health and performance (using tools like Prometheus/Grafana for metrics, ELK stack for logs). In summary, whether cloud or edge, the UKG system is delivered as a containerized application with a focus on reliability, scalability, and portability ³².

Unified 13-Axis Coordinate Schema

At the heart of UKG's data model is the **unified coordinate system** with 13 axes. This schema provides a multi-dimensional "address" for every knowledge element, encoding various facets of context. Each axis has a defined meaning and syntax. Below is an overview of all 13 axes, including their purpose and coordinate notation:

- **Axis 1 – Pillar Level:** Represents the top-level **knowledge domain** or discipline. Pillars are broad fields such as Science, Engineering, Law, Humanities, etc. Every knowledge item is tagged with one (or multiple) Pillars to situate it in a domain of knowledge. *Syntax: 1.<PillarID>[.<SubID>...]* – Pillar IDs are numbered (e.g. 1.32 for Pillar 32). Hierarchical sub-levels can be appended for more specific subdomains or sections ³³ ³⁴. *Example: 1.13.2.2.3* might mean Pillar 13 (e.g. Medicine), then sub-level 2, sub-sub-level 2, sub-sub-sub-level 3, indicating a specific nested topic in that domain ³⁴. Pillar codes often align with major domains (PL01-PL99) – for instance, Pillar 01 might be **Life Sciences**, Pillar 02 **Physical Sciences**, Pillar 32 **Procurement Law** (FAR) ³⁴.
- **Axis 2 – Sector of Industry:** Captures the **industry or economic sector context** of the knowledge. This axis maps to standard industry classification systems (NAICS, SIC, PSC, etc.) and helps contextualize knowledge in a practical domain. *Syntax: 2.<Sector>.<CodeSystem>.<Code>* – for example, 2.6.4 could indicate Sector 6 (say, Healthcare) with code 4 in a certain system ³⁵. A more detailed example: 2.3.3.4 might map to Sector 3 using code system 3 (PSC) with code 4 ³⁵. The system preserves the original codes as metadata (e.g. NAICS 541512 for IT consulting) so that the coordinate 2.x.y.z can be linked back to real-world industry codes ³⁶. Axis 2 ensures that any query can be tuned to the **industry context** (e.g. finance vs healthcare).
- **Axis 3 – Honeycomb (Cross-Domain Links):** A **cross-disciplinary linkage** axis. It represents interconnected nodes that span multiple pillars or sectors, like a honeycomb of knowledge. Axis 3 enables the graph to link a concept into multiple contexts simultaneously. *Syntax: 3.<...>* – often a composite of references to Axis 1 and 2 values, because a honeycomb node links across domains ³⁷ ³⁸. *Example: 3.11.4.2.7* might encode a cross-link connecting Pillar 11 (Health domain), Sector 4 (Information Tech), and some ethical/security concept (the .2.7 part) ³⁸. Essentially, Axis 3 is used whenever knowledge falls in an **interdisciplinary space** – it prevents narrow siloing by formally connecting nodes that belong to multiple areas. (E.g. a concept like "AI in Healthcare Ethics" would be a honeycomb node bridging the medical domain, IT sector, and ethics.)
- **Axis 4 – Branch:** This axis represents **granular sub-classifications or taxonomies within a domain or sector**. If Axis 1 Pillars are broad and Axis 2 is industry, Axis 4 is the fine-grained breakdown inside those. For example, within a Pillar you might have branches for subfields and specialties; within a Sector you might have sub-sectors or categories. *Syntax: 4.<Branch>.<SubBranch>...* (multiple levels as needed). *Example: 4.3.2* could mean Branch 3, Sub-branch 2 – e.g. within an AI domain, Branch 3 might be "Machine Learning" and Sub-branch 2 "Neural Networks" ³⁹ ⁴⁰. This

axis often dovetails with Axis 2 (industry classifications). In regulatory contexts, Axis 4 might capture sub-parts of regulations or programs beyond top-level parts.

- **Axis 5 – Node Identifier:** The **leaf-level node ID** for specific knowledge items. While Axes 1–4 set up hierarchy, Axis 5 pinpoints the exact item (often where a regulation clause, a specific requirement, or a data element resides). *Syntax:* 5.<NodeID>[.<SubNodeID>...] – numeric or alphanumeric identifiers for individual nodes ⁴¹. *Example:* 5.8 could refer to node #8 in a given branch, say a particular clause titled “SOC_ML_Integration” linking social science and ML models ⁴¹. Essentially, Axis 5 is the unique record locator within the context defined by Axes 1–4. Together, **Axes 1–5 define the core hierarchical coordinate** of a knowledge item (Pillar → Sector → Cross-domain → Branch → Node) ⁴².
- **Axis 6 – Octopus Crosswalk:** A cross-reference hub axis for one-to-many links. Octopus nodes are like central reference points that **link one item to many related items** (imagine an octopus with multiple arms). They often represent a **universal concept** that appears in many places. *Syntax:* 6.<OctopusID> ⁴³. *Example:* A concept like “small business” that is defined across multiple regulations could be an Octopus node; 6.5 might be the ID for “Small Business (universal)” which links to all relevant FAR clauses, DFARS clauses, policy docs, etc ⁴³. Axis 6 is heavily used for **regulatory crosswalks** – mapping equivalent concepts across different frameworks so that the system knows, for instance, that “minority-owned business” in one context corresponds to a definition in another law.
- **Axis 7 – Spiderweb Crosswalk:** Another cross-reference axis, but for many-to-many lateral links. Spiderweb nodes create a mesh of connections between analogous nodes in different hierarchies. This is useful for **compliance mappings** where multiple regulations or documents have overlapping provisions. *Syntax:* 7.<SpiderID> ⁴⁴. *Example:* A Spiderweb node might connect FAR Clause 52.219-8 to an equivalent clause in the Small Business Act; both are separate documents, but the Spider node links them ⁴⁴. Axis 7 thus allows modeling of **many-to-many relationships** across the graph (unlike Axis 6 which is one central concept to many references). Between Axes 6 and 7, UKG can represent complex cross-regulatory linkages (Octopus for one concept cited in many places; Spiderweb for mapping several related concepts across laws).
- **Axis 8 – Knowledge Role:** This axis encodes **expert roles or perspectives** associated with knowledge. It answers: *from whose perspective or expertise is this knowledge relevant?* For example, a piece of information might be particularly pertinent to a “Data Scientist” or a “Contracting Officer”. Axis 8 tags such roles. *Syntax:* 8.<RoleID> ⁴⁵. *Example:* 8.12 might stand for the role “Contract Specialist” ⁴⁶. In simulation, this axis is used to route queries to the right **virtual expert persona** – e.g. if a question is legal in nature, Axis 8 = Lawyer persona would activate. (The Quad Persona system uses axes 8–11 to instantiate the four personas, see below.) Multiple roles can be layered via composite codes if needed, but generally one primary role is indicated ⁴⁷.
- **Axis 9 – Qualifications & Skills:** This axis denotes the **credentials or skill level** relevant to the knowledge or required to understand it. It overlays human qualifications (education, certifications, etc.). *Syntax:* 9.<QualID> (sometimes with subcodes) ⁴⁸. *Example:* 9.3.5 could encode Education Level 3 (say, Ph.D.) and Certification 5 (e.g. CPCM – Certified Professional Contract Manager) ⁴⁸. This axis is useful to filter or adapt answers based on the required expertise – e.g. ensuring that if a topic demands a PhD-level understanding, the persona or explanation is chosen

accordingly ⁴⁸. In the knowledge graph, if a node is highly specialized, it might carry an axis 9 tag for the skill needed (like “Board Certified Surgeon” or “AWS Certified Developer” depending on context).

- **Axis 10 – Octopus Regulatory Expert:** A **meta-role axis** representing a composite expert profile that spans multiple regulatory domains. This is one of the special axes used in the Quad Persona simulation. Axis 10 is conceptually an “Octopus Expert” – an agent who has a broad, multi-domain regulatory expertise ⁴⁹. *Syntax:* 10.<ProfileID> (often with an internal structure defining facets of the profile). *Example:* 10.1 might invoke the “Global Compliance Officer” persona profile ⁵⁰. In practice, Axis 10 might not be explicitly present on knowledge nodes; rather, it’s used to **call forth** a persona that can traverse across Pillars. But if a knowledge node is specifically relevant to multi-regulatory expertise, it could be tagged with Axis 10. (In Quad Persona, Axis 10 persona is the *Regulatory Expert* who understands many frameworks.)
- **Axis 11 – Spiderweb Compliance Expert:** Another **meta-role axis** for an expert persona, focused on **harmonizing overlapping regulations**. Axis 11 corresponds to a persona that specializes in compliance integration and conflict resolution across frameworks ⁵¹. *Syntax:* 11.<ProfileID>. *Example:* 11.1 might represent a “Cross-Regulatory Compliance Specialist” ⁵². Again, this is primarily used to spawn the **Compliance Expert persona** in the simulation, who makes sure answers meet all standards and reconciles different rules. In data terms, if a node is specifically about a crosswalk or compliance mechanism, it might carry an Axis 11 tag. But typically Axes 10 and 11 are used in dynamic context rather than as static data properties – they ensure the simulation covers those perspectives.
- **Axis 12 – Location:** The **geospatial or organizational context** axis. It encodes jurisdictional or physical scope for knowledge nodes ⁵³. For regulations, this could be the country or state the rule applies to; for organizational knowledge, it could be the department or unit. *Syntax:* 12.<LocID>[.<SubLoc>...]. *Example:* 12.2.1 might mean *Country = USA* (2) and *Agency = DoD* (1) if, say, 12.2 was USA and .1 under it denotes the Department of Defense ⁵⁴. A different example: 12.5 could be *Region = EU* or *Global = 0* depending on the coding scheme. Axis 12 ensures that knowledge is contextualized to the **correct locality or organizational scope**, so that, for instance, state laws and federal laws can coexist in the graph without confusion ⁵⁵. When a query is asked, if it’s known to be about “California environmental law”, Axis 12 will carry that context so the system filters to California-specific knowledge.
- **Axis 13 – Temporal:** The **time dimension** axis. This anchors knowledge in time, version, or sequence ⁵⁶. It is crucial for managing things like effective dates of regulations, historical data, or step-by-step processes. *Syntax:* 13.<TimeID> (often structured as year.quarter or similar) ⁵⁷. *Example:* 13.2025.1 could denote *Year 2025, Q1* (if .1 is Q1) ⁵⁷, or it could reference a specific event or revision number. The temporal axis allows the system to **handle version control and chronology**, ensuring that answers are based on the laws or data **applicable at the relevant time** ⁵⁷. For example, if a regulation was amended in 2023, the knowledge graph might have one node with 13.2022 (valid up to 2022) and another with 13.2023 for the amended version, each carrying different content.

Each coordinate in UKG is thus a 13-tuple (x_1, x_2, \dots, x_{13}), and not all axes are always used (some may be wildcard or default for a given item). The **Nuremberg-style hierarchical numbering** is used within each axis to show sub-levels (as dot-separated numbers) ⁵⁸. For example, an item with coordinate 1.5.2.1 on

Axis 1 indicates it's in Pillar 1, Chapter 5, Section 2, Paragraph 1 – analogous to a legal outline ⁵⁹. This numbering scheme is human-readable and preserves the **parent-child relationships** between coordinates ⁵⁸. Internally, the system stores these as structured fields and as a composite key, and in the database (USKD) a field like `NurembergNumber` might hold the dotted string for quick reference ⁶⁰.

Finally, every coordinate is paired with **human-readable labels and meta-tags** to integrate with real-world standards. The UKG schema is **compatible with SAM.gov and other official data sources**, meaning when we ingest data, we keep the official identifiers as annotations ⁶¹. For instance, a Pillar that represents the Federal Acquisition Regulation (FAR) will have all its nodes tagged with the actual FAR part/section numbers in metadata ⁶² ⁶³. We might internally refer to a node as `1.32.1.1` (Pillar 32 = Procurement Law, Part 1, Section 1) but the meta-tag will say "FAR 1.101" if that's the corresponding section ⁶² ⁶³. Similarly, Axis 2 nodes carry NAICS or PSC codes as tags (e.g. a node with coordinate `2.54.7` might have metadata "NAICS 5417 – Scientific R&D Services") ⁶⁴ ⁶⁵. This approach ensures no information is lost – users can always map back UKG's coordinates to the familiar references in their domain. The system explicitly integrates with sources like SAM.gov to **ingest official code lists and maintain mappings** ⁶⁶ ⁶⁷. As an example, if we add a new procurement clause from SAM.gov, we import its SAM.gov identifier and attach it to the coordinate.

In summary, the 13-axis coordinate schema is the universal language of the UKG system: it formalizes how knowledge is categorized and linked. This allows the Query Engine and Simulation Engine to reason about "*all relevant dimensions of a problem*" by simply traversing coordinates – e.g. exploring Pillar and Sector for context, Branch and Node for specifics, cross-links for analogies, roles for viewpoints, and location/time for applicability ⁶⁸. All modules of UKG, from data ingestion to query answering, rely on this schema to ensure consistency and coverage.

Core Modules

13-Axis Query Engine (Parsing, Classification & Coordinate Assignment)

The first core module in the reasoning pipeline is the **13-Axis Query Engine**, responsible for understanding user queries and mapping them into the UKG coordinate system. It performs **natural language parsing**, identifies key contextual cues, and produces an initial coordinate (or set of coordinates) that represent the query's meaning in the 13-axis space. This step is crucial because it tells the rest of the system "*what are we dealing with?*" in a structured way.

When a query comes in (e.g. "*An infectious disease specialist needs to design an antibiotic stewardship program in a hospital*"), the Query Engine will:

1. **Parse the Query:** It uses NLP techniques to break down the question – identifying the main subject, the task or problem, any domain-specific terms, any time or location mentions, etc. It may use a combination of keyword analysis, named entity recognition, and possibly a language model embedding to grasp the context.
2. **Classify into Axes:** Based on the parsed information, the engine classifies which **Pillar (Axis 1)** the question falls under, which **Sector (Axis 2)** it pertains to, and so on for other relevant axes. For example, in the above query about antibiotic stewardship, the engine would detect this is a **medical**

scenario (Pillar likely Life Sciences/Medicine) and also related to a **hospital setting** (Sector = Healthcare) ⁶⁹ ⁷⁰. It sees “infectious disease specialist” – indicating a **role context** (perhaps Axis 8 = a medical expert persona) – and “hospital” (Axis 12 could be a location context if specified, or at least an indicator of sector). It also notes the task “design a program to reduce drug resistance” which might imply guidelines or best practices (possibly invoking regulatory knowledge like CDC guidelines – so maybe a regulatory Pillar or cross-domain Axis 3 link to public health).

3. Assign Coordinate Values: The engine then constructs a **coordinate vector** for the query, populating the axes it determined. It essentially translates the natural language into a set of axis values: *Domain, Industry, Location, Time, etc.* ⁷¹. For instance, it might output something like:

4. Axis 1 (Pillar) = Life Sciences (with a subdomain code for Microbiology if it caught that “antibiotic resistance” is microbiology) ⁷²,
5. Axis 2 (Sector) = Healthcare,
6. Axis 8 (Role) = Medical Specialist,
7. Axis 12 (Location) = if the query mentioned a specific region or just “hospital” (could default to general or the user’s org locale),
8. Axis 13 (Time) = Current (or a given year if stated).

In the text, it might explicitly map “2025” to Axis 13 if the question said “in 2025” ⁷⁰. It catches temporal references like “as of 2023” or deadlines. If none given, it may assume current date.

1. Output Initial Coordinates & Context: By doing the above, the Query Engine produces an **initial coordinate or coordinate set that ‘situates’ the problem** in the knowledge space ⁷¹. This coordinate acts as a query vector $\$Q(x)$ for retrieving relevant knowledge. The Query Engine also may output a structured representation of the query (like an internal JSON with fields: domain = Medicine, task = design program, context = hospital, role = infectious disease doctor, etc.). This structured query is essentially the entry point for the simulation: it tells the rest of the system what to focus on.

Under the hood, the Query Engine uses rules and ML models trained on similar queries to classify them. For example, it might have a text classification model for Pillars (to detect if a question is medical vs legal vs technical) and a dictionary of sector keywords (e.g. “hospital” maps to Healthcare sector code) ⁶⁹. It might use a spaCy or NLTK pipeline to extract entities (like dates, locations, roles). All this is then mapped to coordinate codes via a lookup against the schema definitions.

The Query Engine also performs **query validation** – ensuring the query is within scope and not malformed. If it finds the query is gibberish or unrelated to any known Pillar, it can reject or ask for clarification (this is part of Layer 1 in the 10-layer process, often called the Primary Simulation layer) ⁷³ ⁷⁴. This prevents the system from chasing irrelevant paths. For example, an empty query or one with contradictory terms might be caught here.

In summary, the 13-Axis Query Engine is like the “interpreter” that converts a user’s question into the internal language of UKG (the coordinate system). After this step, the system has an **initial coordinate vector representing the query context** ⁷¹, which it will use to fetch relevant knowledge and to initialize the reasoning process. This ensures that from the very start, the simulation is *grounded in the right context* – e.g., a medical question is labeled as medical and will activate medical knowledge, not something unrelated ⁷⁵ ⁷⁰.

(Technical note: In the 10-layer architecture, this Query Engine corresponds roughly to Layer 1 – where the system performs query parsing, axis mapping, and initial memory setup ⁷⁶ ⁷⁷. The output of Layer 1 is a set of coordinates and a “blank canvas” working memory pre-loaded with any obvious context facts about those coordinates ⁷⁸. For instance, if the query is about “antibiotic stewardship”, the system might immediately load a fact that “antibiotic stewardship = program to manage antibiotic use” as preliminary context.)

Simulation Engine (10-Layer Reasoning, Recursive Flow & Refinement)

The Simulation Engine is the core AI reasoning module that takes the structured query from the Query Engine and produces a final answer (with explanations) through a **multi-layer cognitive process**. UKG uses a fixed **10-layer reasoning stack** to simulate the way an expert (or rather, multiple experts) would think through a complex problem ⁵. Each layer has a specific function and the output of one layer feeds into the next. The design is inspired by strategies like “Algorithm of Thought” and “Tree of Thoughts” which break problems into steps ⁷⁹ ⁸⁰, but here the steps are engineered as layers with clearly defined roles.

Layered Reasoning Architecture: The layers can be summarized as follows (conceptually):

- **Layer 1: Query Interpretation & Context Init** – (Handled by Query Engine as described) Parses query, identifies axes, and initializes the working memory with obvious context ⁷⁶ ⁷⁷. Basic sanity checks done here.
- **Layer 2: Knowledge Retrieval (Nested DB)** – Gathers all relevant information from the knowledge graph (USKD) into the working memory. Essentially, this layer says “*Given the query’s coordinates, pull in all the pieces of knowledge that might be useful.*” It queries the in-memory graph and vector indices for nodes that match the domain, keywords, etc. ⁸¹ ⁸². The result is the assembly of a problem-specific subgraph in memory – like fetching relevant laws, data points, prior cases, etc., so that reasoning can happen with that info available ⁸² ⁸³. By the end of Layer 2, the system has a “mini knowledge base” loaded that contains facts and references needed for the next steps ⁸³.
- **Layer 3: Hypothesis Generation (Simulated Research Agents)** – This is where the **Quad Persona** (and potentially more agents) kick in. Layer 3 takes the now understood problem and spawns multiple **specialized reasoning threads** (one per persona and possibly sub-problem) ⁸⁴ ⁸⁵. The idea is to generate hypotheses or intermediate analyses from different perspectives. Each persona agent examines the knowledge from Layer 2 through its lens and starts working on an aspect of the problem. For example, the Knowledge Expert persona will look at factual and scientific aspects, the Industry Expert looks at practical feasibility, etc. They may also identify *sub-questions*: e.g., one agent might say “We need more info on X” and subsequently trigger a retrieval of additional nodes (recursive back to Layer 2 for that sub-query). This layer **divides the problem** among expert agents to ensure thorough exploration ⁸⁴ ⁸⁵.
- **Layer 4: Analysis and Inference** – Each agent (persona) now performs deeper analysis on their sub-problem. This could involve logical inference, applying domain rules, running calculations, or even calling external tools if needed (for example, a math persona might solve equations here). Essentially, Layer 4 is where each thread tries to come up with partial solutions or findings. In the context of Quad Persona, this is each persona’s detailed answer from their viewpoint. The “*Perspective Expansion*” might also occur here – if needed, the system could simulate additional points of view or vary assumptions (like a mini Tree-of-Thoughts for each persona).

- **Layer 5: Synthesis and Collaboration** – The outputs of the persona agents are brought back together. This layer aggregates the findings from Layer 4. The **Aggregation Engine** (mentioned under Quad Persona) collates insights and begins forming a unified answer ⁸⁶ ⁸⁷. If the personas agree, synthesis is straightforward. If there are conflicts (say, the regulatory persona and industry persona suggest slightly different approaches), this layer identifies those conflicts. It may engage in a resolution strategy – e.g., letting the personas discuss (the system can simulate a “debate” or simply apply rules to decide which perspective prevails on each point). By the end of Layer 5, there is a draft answer or at least a consolidated set of points that the final answer should cover.
- **Layer 6: Validation & Cross-checking** – Now the system takes that draft answer and **validates it against the knowledge base and requirements**. This includes checking for factual accuracy (cross-check every statement with the sources loaded in Layer 2), verifying compliance (did the answer consider all regulations? anything violated?), and ensuring consistency. Any gaps identified trigger a backward loop: for instance, if a needed fact is not found, the engine might loop back to Layer 2 or 3 to fetch more info or have an agent address the gap (this is part of the *recursive flow*). Layers 2–6 can iterate in a loop if validation fails: the system refines the answer until it passes checks or reaches a confidence threshold.
- **Layer 7: Optimization & Decision** – In complex scenarios (like planning or optimization problems), this layer would perform any final optimization (e.g., find the most cost-effective solution among the considered ones) or make a decision if multiple answer candidates are present. It applies criteria to choose the best answer formulated.
- **Layer 8: Explanation & Justification** – Here the system compiles the explanation for *why* the answer is what it is. It attaches citations to knowledge nodes used, draws the reasoning chain (so that there is an audit trail). It might produce a structured explanation or bullet points of justification that can later be formatted for the user. Essentially, it ensures the answer is **traceable and justified**, aligning with UKG’s focus on high assurance answers.
- **Layer 9: Final Answer Synthesis** – This layer takes everything (the final content of the answer and the explanation) and synthesizes it into the desired output format (text answer, JSON answer, etc.). It may also format the answer according to user preferences (e.g., if the user wanted a brief answer vs a detailed report).
- **Layer 10: Final Validation (Guardrails)** – The last layer performs any last-minute checks such as ethical guidelines, security (ensuring no sensitive data is accidentally exposed), and format validation. This is like a safety net: if something seems off (maybe the answer contradicts a known law, or the confidence is below threshold), Layer 10 can flag it. It either fixes minor issues or, if severe, it might even halt with a failure message. Assuming all is well, it then outputs the final answer.

This 10-layer process allows a complex query to be handled in a structured way. Importantly, the architecture is **recursive** and iterative – the system can loop back to earlier layers if needed. For example, after Layer 5 synthesis, the system might realize confidence is only 80%. UKG’s design goal is very high confidence (e.g. 99.5%) before output ⁸⁸ ⁸⁹. So it may loop: a *recursive refinement* where it goes back to Layer 3 or 4, asks the personas to consider alternative hypotheses or fill gaps (maybe using different KAs like brainstorming algorithms or counterfactual reasoning). Each pass ideally raises the confidence. This

looping continues until the **confidence threshold is met** or some max iterations occur. This is analogous to an expert reviewing and refining a report multiple times until satisfied.

After the reasoning is done (10 layers and any recursion), the system enters a **Refinement Workflow**. In UKG's design, this is described as a **12-step refinement** that post-processes the compiled answer ⁹⁰. These steps include things like: check logical flow, ensure no contradictory statements, eliminate any biased language, verify all references are cited, polish the language for clarity, etc. ⁹⁰. It's effectively a QA phase on the answer content. Many of these refinement steps are powered by additional Knowledge Algorithms – e.g., a "Logic Auditor" KA, a "Bias Detector" KA, a "Language Clarity" KA. By step 12 of refinement, the answer should be **complete, concise, and correct** (to a very high degree of certainty).

Finally, the Simulation Engine outputs the answer along with metadata like confidence score and the support evidence. If integrated with the UI, it might also output the "perspective breakdown" (what each persona said) for transparency, and the trace of how the answer was developed (for audit logging).

To illustrate, consider the antibiotic stewardship question scenario: - The engine identified Pillar=Medicine, Sector=Healthcare, so in Layer 2 it pulled in knowledge: clinical guidelines from CDC, data on resistance rates, hospital policy templates, regulatory standards (maybe Joint Commission requirements) ⁹¹ ⁹². - In Layer 3, it created persona agents: a Medical Expert persona (to cover scientific aspects), an Operations persona (hospital admin perspective), maybe a Public Health persona (for broader impact), and a Regulatory persona (for any compliance issues). Each persona digs into the Layer 2 info relevant to their view ⁸⁵ ⁹³. - By Layer 5, suppose: the Medical persona suggests a plan emphasizing strict antibiotic usage protocols, the Operations persona adds that staff training and IT systems are needed to track prescriptions, and the Regulatory persona notes any legal mandates (like reporting requirements). They merge this into one comprehensive plan. - Layer 6 checks: does this plan violate any guideline? Did we include metrics to measure reduction in resistance? Maybe it finds no mention of metrics, which the Public Health persona then adds in a recursive pass (say a metric like infection rate). - Eventually, the answer becomes: a multi-point program description which is scientifically sound, operationally feasible, and compliant. The refinement steps ensure it's well-structured and references the guideline documents and data that were used.

The outcome is an answer that *feels like it was produced by a panel of experts deliberating*, because under the hood that's essentially what happened. Thanks to the structured 10-layer approach, the system's reasoning is transparent and each intermediate result can be examined (useful for debugging or audit).

From an implementation standpoint, each layer can be a function or method in code, and the sequence can be managed either procedurally or via a small workflow engine. Knowledge Algorithms (KAs) are mapped to layers (for example, **KA-1: Algorithm of Thought** orchestrates the overall logical flow in layers ⁹⁴, **KA-3: Gap Analysis** might help in Layer 6 to find missing pieces ⁹⁵ ⁹⁶, **KA-12: Role Simulation** runs the personas in Layer 3, etc.). The modular nature means developers can plug in improvements to each layer without affecting others, as long as input/output contracts are maintained.

In sum, the Simulation Engine is the AI's brain that **thinks stepwise** through a problem. The 10-layer architecture ensures complex reasoning is broken down, and the recursive/refinement mechanisms ensure the answer is as accurate and reliable as possible before presenting it. This design is key to UKG's ability to deliver near **error-free answers with traceability**, far beyond a single-pass GPT-style answer generation.

Quad Persona System (Role Instantiation & Memory Flow)

The Quad Persona System is a distinctive feature of UKG's reasoning engine – it allows the AI to simulate multiple expert roles in parallel, giving answers a multi-faceted perspective. Instead of relying on a single reasoning chain, UKG "thinks" as a **team of four experts**, each with a specialized viewpoint, and then combines their insights. This greatly increases the thoroughness and trustworthiness of the results, as each persona can catch issues or add insights that others might miss ⁹⁷ ⁹⁸.

The Four Personas: The core personas in the Quad Persona system are aligned with Axes 8–11 (the role axes) and can be described as: 1. **Knowledge Expert** – Focuses on domain knowledge and factual correctness. (Mapped to Axis 8, e.g. a subject-matter expert role) ⁹⁹. For a given query, this persona ensures the answer is *technically and scientifically accurate*. If the question is about engineering, this is the engineer persona; if medical, the doctor persona, etc. They leverage Pillar knowledge heavily. 2. **Sector (Industry) Expert** – Focuses on practical, real-world context and industry best practices. (Axis 9 role) ⁹⁹. This persona ensures the solution works in practice in the relevant industry. For example, if the question is asking for a business process, the industry persona brings in operational know-how and context-specific constraints. 3. **Regulatory Expert (Octopus)** – A persona specializing in laws, regulations, or overarching policies that span domains. (Axis 10 meta-role) ¹⁰⁰. This expert checks *legal/regulatory implications*, ensuring compliance with any relevant rules. They think broadly across regulatory frameworks (hence "Octopus" reaching into many areas). 4. **Compliance Officer (Spiderweb)** – A persona dedicated to **cross-domain compliance and standards harmonization**. (Axis 11 meta-role) ¹⁰⁰. They focus on consistency and that the answer meets all required standards or policies, especially when multiple frameworks or guidelines are involved. In a way, they're the "safety and policy" checker, often catching conflicts or oversights between different sets of rules.

When a query is processed, **all four personas are instantiated and "activated" in parallel** by the PoV (Point-of-View) Engine ¹⁰¹ ¹⁰². Each persona is essentially an independent agent with its own knowledge filter and reasoning logic: - They **share the initial context** (the relevant knowledge pulled in Layer 2 is accessible to all), but each will zero in on aspects relevant to their role ¹⁰². For instance, given a pile of documents, the Knowledge Expert might gravitate to technical papers, while the Regulatory Expert looks at laws or standards in that pile. - Each persona may also have **role-specific memory or tools**. For example, the Industry Expert persona might have a memory of common industry solutions or case studies; the Compliance persona might have a checklist of compliance standards to verify.

Memory flow and isolation: While the personas operate in parallel, they maintain their own *working memory* during simulation. This means each persona can develop an internal line of reasoning without immediately sharing every intermediate thought. Technically, the system might clone the knowledge graph context for each persona thread (so each can annotate or highlight within that context separately) ⁸⁵. For example, the Knowledge persona might mark certain scientific facts as critical, whereas the Regulatory persona might mark legal clauses as critical. These separate annotated contexts are like each persona's scratch pad.

However, the personas are not fully isolated – the system allows them to **share findings via the Aggregation Engine** as needed. The memory flow is managed as follows: - After each persona has done some reasoning (e.g., completed Layer 4 analysis), the system collects their key outputs (e.g., "Persona A's suggested answer component, Persona B's warnings or additions") into a shared memory or blackboard. - There is an **Aggregation & Synthesis process (Layer 5)** that merges these. At this point, the personas'

individual memory contributions flow into a unified memory structure representing the draft answer ⁸⁶ ⁸⁷. - If a conflict or discrepancy arises, it might trigger a back-and-forth: e.g., Compliance persona flags "Expert's solution violates policy X." This could cause the Expert persona to reconsider or the system to adjust the answer to satisfy compliance. In implementation, this could be another iteration where the conflicting information is fed into the Expert persona's context and Layer 4 is rerun for that persona to address it.

Throughout the simulation, each persona's context is consistently tied to specific axes: - The Knowledge Expert persona "sees" primarily Axis 1 content (domain knowledge) and relevant Axis 3/5 details. - The Industry persona is keyed into Axis 2 and 4 content (sector specifics, branch specifics). - The Regulatory persona leverages Axis 6/10 (reg crosswalks and meta-regulator role) to find broad regulatory principles. - The Compliance persona leverages Axis 7/11 (spiderweb links and compliance role) to check cross-references and standards.

This design ensures **full coverage**: e.g., even if the domain expert forgets a regulatory requirement, the Regulatory persona catches it; if the regulatory person is too strict, the industry persona might counter with practical flexibility.

After they all contribute, their **insights flow into an Aggregation Engine** that synthesizes a unified answer ⁸⁶ ⁸⁷. This is where "the whole is greater than the sum of parts" – the system resolves differences and combines the perspectives. The result is an answer that has, implicitly, been vetted by four experts. UKG documentation describes the end result as having been *informed by all four viewpoints, resolving any conflicts between perspectives* ⁸⁶ ⁸⁷.

Memory-wise, once the personas converge, the separate persona memories effectively merge into one final memory (the final answer context). The persona objects may then be discarded until needed for the next query, though UKG could retain some persona-specific learning (e.g. if over time it notices the Compliance persona often flags a certain issue, it could adjust knowledge base or logic – but that's beyond MVP scope).

The **role instantiation logic** involves mapping a query to these four roles automatically. The Query Engine's classification helps here: if Axis 8 = Scientist, it knows the Knowledge Expert persona should be a Scientist persona; Axis 9 = Finance sector, the Industry persona is a Finance expert; and so on. The system has a library of persona profiles (could be implemented as prompts or parameter sets for the reasoning algorithms) that it loads depending on needed role. For example, *Persona Profile 8.12: Contract Specialist* will have certain knowledge biases and terms it uses. Instantiating that means maybe priming that persona's context with "As a Contract Specialist, I focus on FAR and cost evaluation" etc. The **Model Context Protocol integration** further ensures these roles can be invoked externally as needed (MCP can specify which persona to engage if an external tool only wants one viewpoint, for instance).

In implementation, one can think of each persona as running a modified version of the Simulation Engine focused on a subset of the problem. The coordination is akin to launching 4 asynchronous tasks in Python, each taking the query context and a persona ID, then awaiting all to complete ¹⁰¹. The system then aggregates results. Python's `asyncio.gather` or multi-threading could be used to parallelize this step since each persona reasoning is largely independent computation.

Example of Quad Persona in action: Suppose the query is: "Should our tech company adopt a stricter data encryption standard next year, considering legal and operational factors?" - The **Knowledge persona** (Security

Expert) will reason about encryption standards, cryptographic strength, technical benefits. - The **Industry persona** (Business Ops Expert) will consider operational impact, costs, performance implications in the tech industry context. - The **Regulatory persona** (Privacy/Compliance Lawyer) will bring up laws like GDPR, any upcoming regulations that might require encryption. - The **Compliance persona** (Compliance Officer) will ensure that if multiple standards (ISO, NIST, etc.) are relevant, the recommendation covers them and is feasible compliance-wise. Each works in parallel. The Knowledge persona might say "Yes, we should adopt AES-256 because it's more secure." The Industry persona might caution "This could slow down our systems, we need to budget for better hardware." The Regulatory persona might add "It will likely be mandated by law X next year anyway, so better to comply early." The Compliance persona might unify standards: "Use the standard that meets both ISO and upcoming EU requirements to satisfy all." In aggregation, the answer becomes "Yes, adopt AES-256 encryption next year; while it may introduce some performance overhead (needs hardware upgrades), it will keep us ahead of regulatory requirements (e.g., pending EU law) and align with industry best practices (ISO standard)." This answer clearly reflects input from all four angles.

This system of multiple personas **improves memory coverage** too – because each persona might recall different facts. In UKG's terms, it's like having distinct *memory subspaces* for each perspective that then combine. The memory flow architecture ensures no single persona's limitations become a blind spot for the final answer.

By implementing the Quad Persona System, the UKG MVP achieves a level of robust reasoning similar to a human *expert committee*. The role instantiation logic dynamically brings the right experts to the table for each query, and the memory flow design (parallel processing with synchronized merge) ensures efficiency (it's faster than running four completely separate full simulations sequentially) and effectiveness (broader coverage of knowledge) ⁹⁸. This dramatically increases answer quality and confidence.

Model Context Protocol (MCP) Integration (Context Handoff & Role Activation)

The **Model Context Protocol (MCP)** is an open standard interface for connecting AI models with external tools and applications. In the UKG MVP, MCP is integrated to allow seamless communication between the UKG system and outside clients or automation frameworks, essentially providing a standardized API for the knowledge graph and simulation engine ¹⁰³ ¹⁰⁴. Through MCP, UKG can act as a service that any compatible AI pipeline (like a LangChain agent or a custom enterprise tool) can invoke without needing to know UKG's internal specifics.

Context Handoff: MCP defines a way to package the "context" of an AI interaction (tools, data, constraints) and exchange it. In UKG's case, context handoff means that when an external call comes in via MCP, the relevant information (the query, perhaps user identity or desired persona, etc.) is converted into UKG's internal context (the 13-axis coordinate and initial state), and likewise the results are handed back in a tool-friendly format. The integration works as follows:

- The UKG backend incorporates an **MCP Server/Listener** (this could be part of the FastAPI app or a separate Node.js service using an MCP SDK, depending on implementation choices ¹⁰⁵ ¹⁰⁶). This MCP layer registers a set of "tools" or actions that UKG can perform – for example, an MCP tool named `knowledge_query` or `simulate_scenario`.
- When an external agent calls `simulate_scenario` via MCP, the MCP integration module in UKG will **receive the JSON-RPC (or function-call) request**, which includes the query and any parameters. The module then:

- Maps any provided context into UKG's coordinate system. (For example, if the external call specified a certain domain or provided an existing coordinate, it will align that with UKG's internal axes ¹⁰⁷.)
- Activates the appropriate personas via the Quad Persona system if the call is meant to trigger a full simulation ¹⁰⁸. Essentially, MCP might allow the caller to specify which persona to focus on or to use all four; typically `simulate_scenario` would engage the Quad Persona flow by default.
- Initiates the simulation engine to process the query just as if it came from the UI or API. The difference is, because MCP calls could be part of a larger chain, the integration keeps track of the *session or tool context* (MCP can maintain state across calls if needed, or handle streaming responses).
- As the simulation runs, the integration ensures no information leakage or violation of the MCP contract occurs – e.g., if MCP expects a streaming response, the integration will stream partial results or statuses.
- Once UKG produces an answer, the MCP module packages the output into the MCP response format (could be a JSON object with the answer and confidence, etc.) and returns it to the caller.

One way to think of it: MCP is like a translator between UKG's internal API and the outside world. From the outside, a client might call `UKG.ask(question, persona="Compliance")`; MCP layer takes that, internally calls something like `run_simulation(query, persona_id=11)`, and then returns the result in a standardized format.

Role Activation via MCP: MCP also provides a way to externally trigger specific **persona roles or tools** within UKG. For example, an external orchestration might decide: first ask UKG's "Regulatory Expert" persona a question, then ask the "Industry Expert" persona, then combine answers itself. To support this, the MCP integration exposes perhaps tools like `ask_regulatory` or a parameter on `simulate_scenario` to only use a certain persona. When such a call is made, the integration will **activate only that persona's reasoning** internally (perhaps by restricting the Quad Persona system to just that axis). In effect, this allows external systems to use UKG in a modular way – either get the combined answer or query individual perspectives.

From a development standpoint, integrating MCP meant implementing handlers that map **MCP's JSON-RPC interface to UKG's Python functions**. The design ensures that all of UKG's unique features – like the 13-axis context, the recursion loop, the persona orchestration – are *encapsulated behind a single MCP call* if desired ¹⁰⁹. For instance, an external caller can invoke `simulate_scenario` once and **UKG will internally do axis mapping, run multi-axis reasoning, recursively refine until confidence ≥99.5%, and return the final result** ¹⁰⁹ ¹⁰⁷. This makes UKG's complexity "disappear" behind the MCP interface – clients don't need to manage the multi-step process, they get a straightforward call that yields a high-quality answer.

The integration also handles **long-running interactions**: If a simulation might take longer, the MCP spec allows either streaming intermediate results or using an asynchronous pattern (acknowledge the request and later send the result). The UKG MCP module is configured to handle synchronous queries quickly (most should finish in <1s in memory) ¹¹ ¹¹⁰. For very complex cases, it could either stream a status (like "still thinking...") or the deployment can opt to raise timeouts.

Additionally, through MCP, UKG can **expose fine-grained tools**: e.g., an `lookup_coordinate` tool to fetch info on a given coordinate, or a `list_personas` tool to list available personas. This turns UKG into a

versatile service accessible to other AI systems (like you could plug UKG as a knowledge provider into a broader agent that also uses other tools).

The **benefit of MCP integration** is that it makes UKG interoperable. For example, an enterprise might use LangChain or Microsoft's Orchestration to chain a question to multiple systems – with MCP, UKG can be one of those systems easily invoked ¹¹¹. It's called the "*USB-C of AI apps*" in concept ¹¹² – a universal connector. By aligning with this standard, UKG can be dropped into many infrastructures without custom integration code.

To configure this, developers ensure the MCP server (whether in Python or Node) is properly authenticated and secure (since it could be an entry point to powerful UKG functions). In MVP, the MCP interface might be optional, but providing it sets the stage for easy extension.

In summary, the MCP module makes UKG's advanced capabilities available as a **standardized AI service**. It translates external requests into UKG's coordinate-driven simulation calls and then returns the results in a tool-friendly format ¹¹³. It handles activating the correct personas and mapping the multi-axis context during this handoff. This allows an engineering team to, for instance, plug UKG into a chatbot or a workflow system by simply making MCP calls, rather than dealing with the complexity of coordinate systems and persona logic themselves. The integration extends UKG's reach and makes it a component in a larger AI ecosystem, following best practices for AI interoperability.

Unified Coordinate System Logic (Nuremberg Numbering & SAM.gov Integration)

The **unified coordinate system logic** underpins how every knowledge node is identified, stored, and linked. Two key aspects of this logic are the use of **Nuremberg-style hierarchical numbering** for structuring coordinates and the **integration of external standards (like SAM.gov codes)** to enrich and align coordinates with real-world references.

Nuremberg Numbering for Hierarchy: The coordinate system uses a dot notation reminiscent of legal document outlines (often called Nuremberg numbering) to represent hierarchy on each axis ⁵⁸. This means that a coordinate isn't just a flat number – segments separated by dots indicate parent-child relationships. For example: - A Pillar coordinate **1.5.2.1** can be interpreted as Pillar 1, Section 5, Sub-section 2, Clause 1 (if we analogize to a document) ⁵⁹. Pillar 1 is the parent, 5 is a child under it, 2 under that, etc. - A Branch coordinate **4.10.3** might mean Branch 10 with a sub-branch 3 (so 10 is parent category, 3 is a subcategory).

This numbering scheme is **human-readable and maintainable**. It allows engineers and domain experts to navigate the knowledge graph more intuitively. For instance, if two nodes have coordinates **1.5.2.1** and **1.5.2.2**, one can tell they are sibling nodes under the same context (maybe two clauses in the same section) without needing to query the database. Internally, UKG stores these as either strings or as separate fields. In the USKD schema, a table like **Regulations** might have a column **NurembergNumber** to store the dot-separated outline number for each clause ⁶⁰. This ensures traceability: if someone has a reference "Section 5.2.1 of the Procurement Law", it corresponds to coordinate **1.5.2.1** which is directly stored.

The logic includes functions to **parse and compare hierarchical coordinates**. For example, given **1.5.2.1** and **1.5**, the system can infer that the former is under the latter (child of it). Or given **2.16.4** and **2.16.10**, it knows they share a parent **2.16**. This is heavily used when traversing the graph: if a

query is at a high level, the engine might iterate through sub-nodes by incrementing those numbers logically.

The numbering follows each axis's own taxonomy depth. Some axes might have fixed depth (e.g., Axis 13 Time might always be Year.Month.Day or Year.Quarter), others are variable (Pillars can have arbitrary subdepth for nested subsections). The coordinate logic is flexible to handle this – e.g., store coordinates as arrays of integers for each axis.

From a development perspective, utility functions like `isAncestor(coordA, coordB)` can be implemented to check if coordA is a prefix of coordB in dot terms (and axes match). Another function `nextSibling(coord)` might generate the next coordinate in sequence at the same level (if needed for traversal). The numbering scheme also helps when sorting nodes – lexicographically sorting the dot numbers usually corresponds to the logical order in a document.

SAM.gov and External Code Integration: The UKG coordinate logic is augmented with **metadata linking to external standards**, particularly from government data sources like SAM.gov (System for Award Management) which provides official industry and procurement classification codes. The idea is that while UKG internally uses numeric codes for axes (for uniformity and performance), it never loses the original context of those codes.

For example: - If Axis 2 coordinate is `2.54.7`, UKG might know internally that "54" corresponds to NAICS Sector 54 (Professional, Scientific, and Technical Services) and "7" might correspond to a specific industry code (like 5417 which is Scientific R&D Services) ⁶⁴ ⁶⁵. The coordinate is stored as numbers for quick computation, but alongside it, metadata stores "NAICS 5417" and the textual description. - If Axis 1 has a Pillar that is actually the FAR (Federal Acquisition Regulation) as Pillar 32, UKG doesn't just say Pillar 32 – it tags that node with `FAR` so humans know it's FAR. Thus a full coordinate might be displayed as `FAR.1.1.1` meaning actually Pillar 32 internally ⁶² ⁶³.

The system is designed to be **fully compatible with naming standards**: whenever possible, a node's metadata includes the exact identifiers from its source domain ⁶¹. If we ingest Federal regulations, each section keeps its CFR or FAR numbering in a field. If we ingest military standards, their MIL STD numbers are kept.

This has two big advantages: 1. **Traceability:** Users can cross-reference any UKG output with real-world references. For instance, if UKG answers "according to FAR 9.1, you must do X", the coordinate might be something like `Pillar 32 node 9.1` internally, but because we stored "FAR Part 9.1" as a meta-tag, the system can output the familiar citation ⁶² ⁶³. 2. **Integration:** If UKG needs to integrate with external databases or APIs (say fetch data from SAM.gov or verify a code), having the official code handy means it can directly query those external sources by code. For example, if a new NAICS code appears, UKG can automatically align it because it uses the same coding scheme.

The logic for SAM.gov integration in practice involved loading the standard code lists into UKG's knowledge base. For Axis 2, UKG likely loaded all NAICS codes, SIC codes, PSC codes, etc., each as nodes or part of the schema. Each code node has both a numeric coordinate and the official alphanumeric code as a tag ⁶⁴ ¹¹⁴. The coordinate might condense an alphanumeric code into numeric form for uniformity (e.g., "R425" PSC code could be mapped to some 3-part numeric code), but in metadata it will say "PSC R425: Support-

Professional: Engineering/Technical". So when that appears in an answer or interface, users see the real code.

For government contracting, SAM.gov data includes things like responsibility codes, exclusion records, etc. UKG could incorporate some of these as well – the main point for MVP is classification codes and standard entity identifiers.

The coordinate system logic likely includes a **mapping table or dictionary** for each axis that corresponds to an external standard: - Axis 1: mapping of Pillar numbers to domain names (and acronyms if applicable, like mapping 32 → "Procurement Law (FAR)"). - Axis 2: mapping of Sector and code system to actual codes (like 2.6.4 → NAICS 511210 if that was how it's structured). Possibly UKG uses subfields in Axis 2 coordinate for "code system" and "code", enabling multiple systems. - Axis 8-11 (roles): mapping role IDs to titles (e.g., 8.12 → "Contracting Officer (GS-13)" as given in an example ¹¹⁵). - Axis 12: mapping location IDs to names (1 = Global, 2 = USA, 2.1 = USA->DoD, 3 = EU, etc.). - Axis 13: mapping some temporal markers if needed (though time can often be self-evident like year numbers).

UKG likely provides an API or function to **translate a coordinate to a human-readable string** using all these meta-tags. For instance, a knowledge node might be represented as "*Federal Acquisition System (FAR 1.102) [FAR.1.1.1.1.1]*" where the part in parentheses comes from metadata and the part in brackets is the full coordinate ¹¹⁶ ¹¹⁷. This way, both humans and machines are happy: machines use the bracketed coordinate, humans see the familiar names.

One more facet: the logic ensures **bidirectional lookup** – given a citation like "FAR 9.1" one can find the UKG node, and given a UKG coordinate one can find the original citation ¹¹⁸ ¹¹⁹. This was explicitly a goal in the design: "*bidirectional lookup: given a UKG coordinate one can retrieve the original citation (and vice versa)*" ¹¹⁸. Thus, developer guidelines include maintaining dictionaries for acronym-to-Pillar and code-to-Axis2 mappings, etc.

Finally, because UKG aims to be comprehensive, its coordinate logic is extensible. If a new axis or coding system comes into play, it can add that while preserving integration. For example, if in future an *Axis 14* (hypothetical) for "Energy Efficiency Score" was needed, the system's modular design (see extension strategy below) could incorporate it ¹²⁰ ¹²¹.

In essence, the Unified Coordinate System logic ensures every piece of knowledge is **uniquely identified, hierarchical, and enriched with real-world semantics**. Developers working with the system should always assign coordinates to new data following the established syntax, and attach meta-tags for external references. The framework already includes functions to generate new coordinates (e.g., when adding a new node under a branch, it will find the next available number in that hierarchy) and to validate coordinates (ensuring they conform to expected patterns like no skipping of levels without an intermediate node).

By adhering to these conventions, the UKG MVP maintains a rigorous structure where knowledge can be expanded and integrated without chaos – everything stays "in place" in the 13-dimensional library, and anyone can find it by either technical coordinate or common name. This makes the knowledge base scalable and interoperable with external systems, which is especially important in enterprise and government contexts where standard codes are the lingua franca for data sharing.

Data Layer (Knowledge Graph and Storage)

The data layer of UKG encompasses the **Universal Simulated Knowledge Database (USKD)** – essentially the knowledge graph – along with any supporting storage systems like vector indexes and external databases. The MVP focuses on an in-memory knowledge graph for rapid simulation, complemented by persistent storage for larger scale or durability, and uses modern data technologies to handle both structured relations and unstructured semantic search.

In-Memory Knowledge Graph (USKD): The primary data structure is an in-memory graph that holds nodes (entities, facts, clauses, etc.) and edges (relationships). The simulation engine operates directly on this in-memory graph to achieve high speed (no DB lookup latency during reasoning) ⁷. The graph is likely implemented using a library like **NetworkX** (for Python), which allows complex graph queries and manipulations in memory ¹²². Each **node** in the graph has: - A unique 13-axis coordinate (as discussed) which serves as its primary key or composite key. - A set of properties/attributes, including a human-readable name/title, a body of text or value (for example, the text of a regulation clause, or a data point value), and metadata (source, effective date, etc.). - Vector embeddings for semantic content: if the node contains text (like a paragraph of text or description), the system may store a high-dimensional vector embedding of that text. This is used for semantic similarity search in the vector store (more on that below). - Links (edges) to other nodes representing relationships. Edges can have types, e.g. “is part of”, “refers to”, “cites”, “contradicts”, etc. In the regulatory domain, an edge might connect a FAR clause node to a corresponding DFARS clause node (that could be a Spiderweb relationship in axes but also implemented as a graph edge for quick traversal).

The **graph database integration** likely uses a hybrid approach: - For persistence and complex querying, a **Neo4j** or similar property graph database is used in the background ¹²². Neo4j excels at querying path relationships (like “find all nodes connected to X via a certain relation”) and can store billions of nodes if needed. In MVP, the data volume might be small enough that not all of Neo4j’s capabilities are needed, but the design anticipates scaling. The system can push updates to Neo4j for storage, and on startup load relevant subgraphs into memory (or even the entire graph if it’s small enough). - For the active simulation, a **NetworkX in-memory graph** is maintained, possibly limited to the portion of the graph needed for current operations (like the subgraph relevant to the query) ¹²². However, since MVP might not be memory-constrained with moderate data, it could load everything into memory. This “RAM-resident” approach is emphasized to ensure the simulation is traceable and self-contained ⁷ – no unpredictable delays from external DB calls mid-reasoning.

UKG also mentions dynamic graph behaviors like **node cloning** and **graph partitioning** for optimization ¹²³. For instance, if new info arrives that closely resembles an existing node (by cosine similarity), the system might clone that node’s structure to quickly integrate it (this is an advanced feature to keep graph topology optimal) ¹²³. Clustering (e.g., “JAEGER partitioning with k-means”) is used to segment the graph by topics for faster retrieval ¹²⁴. These are algorithmic optimizations; for MVP, the key is that the data layer supports such operations – meaning we have the ability to compute similarity between nodes (hence we store embeddings), and we can label clusters of nodes.

Vector Store Integration: In addition to the graph’s symbolic relationships, UKG uses a **vector index** to enable semantic searches. This is crucial for the Knowledge Retrieval step (Layer 2) where it might not be obvious which nodes are relevant via direct links; the system can do an embedding-based search. The MVP might integrate with a vector database like **Milvus** or **FAISS** (Facebook AI Similarity Search) for this purpose.

For example, all textual nodes have embeddings stored in Milvus, so when the query is embedded, the system can ask Milvus for the nearest neighbor nodes (which might find relevant pieces of text even if not directly linked by an explicit relationship) ⁸¹ ¹²⁵.

The integration likely works as: - When new data is ingested, generate an embedding (using a model like BERT or a domain-specific transformer) for the text and upsert it into the vector store with the node's ID. - At query time, the Query Engine also embeds the query (or key terms from it) and queries the vector store for top-\$k\$ similar items. The results come back as a list of node IDs with similarity scores. - These nodes are then pulled into the working memory graph for consideration (that was described in Layer 2 retrieval) ⁸¹ ⁸². - Because all nodes have coordinates, even if we found something via semantic similarity, we know its coordinate and can thus place it in the correct context as we integrate it.

Milvus, for instance, can handle large vector collections efficiently and allows filtering by metadata as well. UKG could use that to, say, only search within a Pillar subset if Pillar is known, improving relevance.

Data Schema for Nodes: Likely, the knowledge is divided into categories/tables such as Regulations, DomainKnowledge, CompanyData, etc., but since it's a graph, each node can be different but still have some common fields. A pseudo-database schema might be:

```
Nodes(node_id PRIMARY KEY, coordinate TEXT, title TEXT, content TEXT, embedding  
VECTOR, metadata JSONB)  
Edges(edge_id PRIMARY KEY, source_node, target_node, type TEXT)
```

If using Neo4j, we'd label nodes by type (e.g., RegulationNode vs PersonNode if people included, etc.) and use relationships.

One specific example provided: a **Regulations** table with fields like regulation_id (UUID), regulation_type, x_coordinate, y_coordinate, z_coordinate, w_coordinate (like a 4D subset maybe used earlier), etc. in the Mathematical Framework doc ¹²⁶. That earlier 4D might have been a simplified indexing (like Domain, Hierarchy, Relationship, Role dims as an example). The MVP's 13D is more detailed.

The **format for knowledge node storage** in practice: - The **coordinate mapping** is either stored as multiple columns (one for each axis) or as one string. Storing each axis separately (like 13 columns) could allow indexing and numeric comparison (for hierarchy), but storing a single string is simpler. Perhaps both are done: one composite string and also separate fields for major axes like Pillar, Sector, etc. (This is speculation, an implementation detail choice.) - **Metadata** is stored in a JSON or properties. This includes things like official names, source references, authors, etc. The meta-tag conventions described earlier ensure important identifiers (like "FAR Part X") are present ¹²⁷ ¹²⁸. - **Vector Embeddings:** if using a Postgres, they might use the PGVector extension to store the embedding in the same DB; if using a specialized vector DB (Milvus), it stores outside but references the node_id.

The data layer also must support **updates**: If new knowledge is added, the system will assign it a coordinate (perhaps the next number in the relevant branch) and insert the node in both the graph and vector index. If knowledge is updated or a regulation changes (versioning), the old node might get a "archived" flag or a new temporal axis value, and a new node added with updated content (Axis 13 handling). The relational

integrity is maintained by linking the new node to the old (maybe an “amends” edge) and using Axis 13 to denote the timeline.

For **graph database integration**, if using Neo4j: - UKG could mirror key parts of the in-memory graph to Neo4j so that if the system restarts or for analytical queries you can run Cypher queries on the full graph. For example, one could query Neo4j: “find all regulation nodes related to cybersecurity that are connected to healthcare sector nodes” to identify cross-domain regs. This is outside the direct simulation, but useful for knowledge management. - The MVP could initially avoid needing Neo4j by storing things in-memory or in simpler stores, but given future scale, code structure should allow plugging in Neo4j queries where appropriate.

Integration Example: Suppose UKG ingests a new policy document. The ingestion pipeline (which might be offline or separate) parses it, assigns coordinates to each section, creates nodes for them (with text content), links them to the Pillar and Sector, tags them with any known codes (like if it's a DoD policy, link to DoD node on Axis 12, etc.), computes embeddings for each section, and adds to vector index. Now the next time a query comes that these nodes are relevant for, Layer 2 can find them either via direct coordinate traversal or semantic search, and they'll be present for the AI to use.

On the **vector store side**, if using Milvus, one could create collections like “RegulationEmbeddings”, “ResearchPaperEmbeddings” etc., or a single unified collection with tag filters. Query Engine might then do something like: `query_embedding = model.embed(query); results = milvus.search(collection="AllEmbeddings", vector=query_embedding, filter={"pillar": [pillars_detected]}, top_k=20)`. This returns nodes IDs which the knowledge service then retrieves (from memory graph or if not present, from persistent store).

Performance considerations: The in-memory approach gives very fast read/write during simulation, but if data is huge (millions of nodes), memory might be an issue. MVP probably deals with manageable data (maybe thousands of nodes like all relevant regulations and some internal data). A cloud deployment might still run on a single beefy instance in memory. For scale-out, approaches like partitioning by Pillar (some distributed graph) would be considered, but MVP likely doesn't require that complexity.

In summary, the data layer combines the **strengths of graph databases and vector databases**: - Graph structure for explicit relations (ensuring context and traceability). - Vector similarity for implicit relations (ensuring recall of relevant info). - In-memory operation for speed, with persistent backing for data integrity and volume beyond RAM if needed ¹²². Developers interacting with this layer will typically use high-level functions (like `get_related_nodes(coord)` or `search_by_text(query)`) rather than raw queries, as the complexity is under the hood. The documentation for MVP should include instructions on how to add new data (e.g., “to add a document, assign coordinates to its parts using the schema, then call the `graph.add_node()` and `vector_index.add()` functions accordingly”). This ensures the knowledge graph can grow and evolve while maintaining the unified structure.

UI/UX Design and Component Library

The UKG MVP's user interface follows a cohesive UX design, with a focus on clarity, ease of navigation, and compliance with accessibility standards. It uses a **custom component library** guided by the design system

(as touched on in the frontend architecture), which ensures consistency in look and feel across the app's various features. Below, we detail key design and UX elements:

- **Theming and Modes:** The UI supports both **Light and Dark modes**, allowing users to switch based on preference or environment. In light mode, the background is generally white or light gray, and text is dark; in dark mode, backgrounds are dark (e.g., charcoal #1a1a1a) and text is light ¹⁷. The design uses a set of **design tokens** for colors, typography, and spacing so that themes can be applied uniformly. For example, the primary color token is deep blue (#003366) and secondary is teal (#008080) ¹⁵ – these are used for headers, accents, and buttons. A change in theme might adjust these slightly (e.g., in dark mode, primary color might be used more for text or outlines rather than large fills to reduce glare). All components reference these tokens (for instance, buttons use primary color for background in the primary style). Spacing tokens define things like base unit (4px or 8px increments), container padding (perhaps 16px), etc., to keep layout consistent.
- **Layout and Navigation:** The app layout is divided into a **sidebar navigation**, a **top navigation bar**, and the **main content area**.
- The **Top Nav Bar** is a fixed header (e.g., 64px tall) that spans the top of the screen ¹⁸. It contains the application logo on the left, the app title, and on the right side, user-related controls (profile menu, language selector, theme toggle, etc.) ¹²⁹. It may also include quick-access icons (like a notification bell or help icon) as needed. The header uses the primary brand color as background in light mode (or a dark variant in dark mode) to clearly delineate it.
- The **Sidebar** is a vertical menu on the left, typically 240–280px wide, which lists the main sections of the application (for example: Dashboard, Knowledge Base, Queries, Reports, Admin, etc. depending on features). It is collapsible – when collapsed, it might show only icons, and on hover or click it can expand to show labels ¹⁹. The sidebar in dark theme uses a dark background (#1a1a1a) ¹³⁰, and in light theme likely a lighter color or white. Menu items are styled with a highlight for the active section (e.g., a blue line or background on the selected item). There's also a **search bar** at the top or bottom of the sidebar for quick navigation or finding content ¹³¹ (like a quick jump to a knowledge node by keyword).
- **Breadcrumbs:** At the top of the main content area (just below the top nav bar), a breadcrumb trail is displayed when the user is in a subsection of the app or viewing a specific knowledge node. For example, if the user navigated to *Knowledge Base > Regulations > FAR > Part 9 > Section 9.1*, the breadcrumbs would reflect that path. This provides context and one-click access to higher-level pages ²⁰.
- **Main Content Area:** This is a scrollable container that holds the actual page content. It uses a **responsive grid** (12-column layout) to arrange content such as cards, tables, or forms ²¹. On wide screens, multiple columns of cards can display, whereas on mobile, it will collapse to a single column. A consistent gutter width (24px) is used between columns ²¹. The content area may have different layouts depending on the page (e.g., a dashboard page might use a card grid, whereas a detailed view page might use a 2-column layout with content and a sidebar of metadata).
- **Component Library:** The app's components are built to be reusable and styled consistently. Some key components include:

- **Cards:** Used for presenting a summary of an item (knowledge node, search result, statistic, etc.). The design for cards typically is a white (light mode) or very dark gray (dark mode) background with a subtle drop shadow (e.g., `0 4px 6px rgba(0,0,0,0.1)`) and rounded corners (8px radius) ²². The card component often includes a header (maybe an icon or title) and body text or chart. Cards are sized to a uniform width (e.g., 300-360px) for grid alignment, with flexible height depending on content ²². They also have defined states: a hover state (perhaps slightly elevate or brighten the shadow), a selected state (outline or highlight).
- **Buttons:** There are primary and secondary button styles. Primary buttons use the brand primary color (#003366) as background with white text ¹³²; secondary buttons might have a transparent or light background with a teal border and text ¹³³. Hover effects cause a scale-up (105%) with a 0.3s transition and maybe a slight shade change ¹³⁴. Disabled buttons are faded and not interactive. Button sizes are standardized (e.g., 40px height for normal buttons with 16px horizontal padding) ¹³⁵.
- **Forms and Inputs:** Input fields (text inputs, textareas, dropdowns) are styled with a consistent height (around 48px for text inputs) ¹³⁶, padding (12-16px), and border radius (maybe 4px or more if rounded style) ¹³⁷. They use neutral border colors (#ccc) which turn into primary color on focus to indicate selection. Each input has an associated label; we ensure every label is properly linked to its input for accessibility (using `<label for="id">`). Error states are indicated by a red border or message below the field. For multi-select or typeahead inputs, the dropdown menus follow the same style as cards (light background, shadow).
- **Modal Dialogs:** A modal appears centered on screen with an overlay backdrop. It uses the same card styling for consistency – usually with a slightly larger padding. The modal header might have a bold title, and the modal content can contain text or form elements. The primary action (e.g., “Save”, “Confirm”) is a button on the bottom right, with secondary (“Cancel”) on the bottom left.
- **Navigation Components:** Aside from the main nav, sub-tabs or secondary menus within a page are styled as tab buttons or pills. If, say, within the Knowledge Base section there are tabs for “Regulations”, “Policies”, “Guidelines”, those tabs are perhaps horizontal at top of content area.
- **Data Visualization Components:** If showing charts (like analytics dashboards), the color palette for charts is derived from the theme (using brand colors and complementary ones). We ensure charts have proper legends and tooltips (with accessible text).
- **Knowledge Graph Viewer:** A specialized component (if included) might show a network graph of related knowledge nodes. This could be implemented with D3 or Three.js as mentioned ¹³⁸ ²⁷. It would allow zooming/panning and clicking nodes to see details. Visually, nodes could be color-coded by Pillar or Sector, using the design system colors.
- **Accessibility Features:** The UI is built to meet **WCAG 2.1 AA** guidelines:
 - Contrast: Text and UI elements have sufficient contrast against backgrounds. For instance, the deep blue (#003366) on white is good for headers, and for dark mode, light text is used on dark backgrounds. We also provide a “high contrast” toggle (especially for dark mode) which may increase contrast further for users with visual impairments ²⁹.
 - **ARIA Labels:** All interactive controls (buttons, links, form fields) have descriptive labels or `aria-label` attributes for screen readers ¹³⁹. Icons (if any) either have accompanying text or are labeled. For example, the search icon button might have `aria-label="Search"`.
 - **Keyboard Navigation:** The interface can be fully operated with a keyboard. The focus order is logical (e.g., top bar items first, then sidebar, then main content). Skip links (like a hidden “Skip to main

content” link) are provided for keyboard users to jump directly to content, bypassing repetitive nav links ¹⁴⁰. Focus styles are clearly visible (e.g., a thick outline on focused buttons or links).

- Responsive Design for assistive tech: The layout can adapt not just to screen size but to different zoom levels or if large text is enabled on the OS. The use of relative units (em, rem) for fonts and spacing helps in this regard.
- Testing: The team would test with screen reader software (NVDA/JAWS) to ensure the reading order and announcements make sense. Also test for operations like closing modals with Esc, navigating menus with arrow keys, etc.
- **Design Tokens and Documentation:** The design system is documented so developers know how to use it. For example:
 - Typography scale: perhaps H1 = 48px bold, H2 = 36px bold, H3 = 24px bold, Body = 16px regular, etc. ¹⁴¹ ¹⁶.
 - Color usage: primary for headers and accents, secondary for highlights or links, success (green), warning (orange), danger (red) colors defined for status indicators.
 - Spacing scale: e.g., 8px base grid, common spacing values = 8, 16, 24, 32, 48px etc. Section margins at 48px, etc. ¹⁴².
 - Iconography: using a consistent icon set (Material Icons or FontAwesome) with line style matching the design (perhaps outlined icons in dark mode, filled in light).
 - Interaction states: documented how components react on hover, active, disabled, focus states.

Given this robust UI framework, when implementing new features, engineers have a toolbox of components and styles to use, ensuring the MVP remains visually consistent and user-friendly. The UI/UX design ultimately aims to make the powerful capabilities of UKG accessible: users can easily navigate the knowledge graph, input queries, and understand the answers (with clear formatting and the ability to drill down into sources, likely by clicking citations on answers to see the underlying knowledge nodes in modals or side panels).

Deployment Models

UKG MVP is designed to be deployable in various environments to meet different requirements. Three primary deployment models are considered: **cloud-first**, **edge/offline**, and **hybrid cloud-edge**. Each model uses the same core software components but packaged or configured differently for its context.

- **Cloud-First Deployment:** This is the default deployment model, where UKG runs on cloud infrastructure (public or private cloud). In this setup, all components (the FastAPI backend, databases, vector store, front-end) are hosted on cloud servers or Kubernetes clusters. A typical cloud deployment might look like:
 - A Kubernetes cluster running the UKG backend as a set of containers (with auto-scaling enabled). The cluster might have multiple replicas of the FastAPI app behind a load balancer to handle concurrent queries.
 - Managed cloud services for the databases: e.g., a Neo4j Aura instance (Neo4j’s cloud service) for the graph database, a managed PostgreSQL or Milvus service for vector embeddings, etc. These can also be containerized if using an on-prem k8s cluster.

- The front-end web app served via a CDN or cloud storage bucket for static content, or via a container (nginx or node server) in the cluster.
- **Networking & Security:** All communication is over HTTPS (with TLS certificates managed via something like Let's Encrypt or cloud certificate manager). The FastAPI might sit behind an API gateway or an Ingress that handles TLS termination. Authentication in cloud model might integrate with cloud identity providers (OAuth with Google, Azure AD, etc., if enterprise).
- **Scaling:** In cloud, if usage grows, the stateless parts (app servers) scale horizontally. The knowledge graph DB and vector DB might need vertical scaling or sharding – e.g., Neo4j cluster or Milvus cluster scaling out. But MVP likely runs on a single instance of each with enough resources.
- **CI/CD:** A cloud-first approach assumes frequent updates, so a CI/CD pipeline is set up. On merge to main branch, pipeline runs tests then deploys new container images to the cluster (possibly rolling updates). This ensures the service is always up-to-date. Cloud monitoring is set (Prometheus, CloudWatch, etc.) to track performance and any issues.

Benefits: Centralized knowledge base, easy to manage, all users get updates immediately. Suited for enterprise deployments where central IT manages the service and users access via internet or VPN. Data can be kept secure with proper cloud security (VPC isolation, etc.). The cloud model is referenced in design as the primary approach (FastAPI + container stack) ³².

- **Edge/Offline Deployment:** This model packages UKG to run on-premises or in an environment with limited or no connectivity to the central cloud. This is useful for clients who have sensitive data that cannot leave a local network, or for scenarios like a field operation with no internet.
- In an edge deployment, **UKG is delivered as a self-contained package**, for example, a set of Docker containers or even an installable VM/ appliance. It includes the core knowledge graph data (or a relevant subset) pre-loaded. The phrase "*pre-packaged knowledge nodes*" means the necessary chunk of the knowledge base is exported and baked into the local deployment image. Possibly, a static snapshot of the USKD (like a database file or dump) is shipped and loaded on startup.
- Because edge devices might not be as powerful as cloud servers, optimizations are considered: for instance, smaller language models or no heavy neural dependencies unless a local GPU is available. The MVP might allow toggling certain features – e.g., if no vector DB installed at edge, rely on pre-computed links more.
- **GPU Acceleration:** If the simulation uses any heavy ML (like embeddings or certain KAs), having a GPU on the edge device can drastically speed it up. So the edge model would be configured to use local CUDA for any ML model inference. Alternatively, some models could be distilled to run on CPU if GPU not available, trading off speed.
- The user interface is the same, but served from the local instance. Users connect to a local URL or IP. No external calls need to be made (all data and processing is on device).
- **Deployment footprint:** Could run on a single server or even a beefy laptop. Docker Compose or Kubernetes (k3s for lightweight) could orchestrate multiple containers (one for backend, one for DB, etc.) on that machine. We ensure the package includes all dependencies so it can run offline (for example, if any third-party data was fetched during simulation in cloud, we'd need to mirror or cache it for offline use).
- Edge deployments also consider data updates: the knowledge graph on edge might need periodic updates (e.g., if new regulations come out). Without connectivity, one approach is to ship update packages (maybe USB or file transfer) periodically that can be applied to the edge instance to import new nodes.

Use case example: A defense contractor runs UKG on a secure network without internet. They get a quarterly update file from the central team with new FAR/DFARS changes, which they import into their edge UKG. They can then query it internally with no data leaving their network.

- **Hybrid Cloud-Edge Deployment:** In this model, you have both a central cloud instance and one or more edge instances, working in tandem. The hybrid approach aims to combine the strengths of both: the cloud does heavy lifting and holds the global data, while edge provides low-latency local access and operates even when offline, syncing when possible.
- **Cloud-Edge Sync:** The core mechanism here is a synchronization service. The cloud might host the master knowledge graph; edge nodes subscribe to updates relevant to them. For example, an edge instance at a regional office might only need certain Pillars or data related to their operations. The cloud can push updates of those parts to that edge. Sync could be via secure REST calls, message queues, or even periodic pull (edge calls cloud for diff). Conflict resolution strategies are in place if edge can modify data (though likely edge is mostly read-only or adds local-only data).
- **Query Federation:** In some hybrid setups, an edge instance might forward queries it can't handle to the cloud. E.g., if an edge has only unclassified data but a query asks something that requires classified data in cloud, maybe it will call the cloud. But often, edge will try to handle everything locally to minimize dependency, using cloud only for updates or exceptionally large tasks.
- **Example:** Consider an oil company with oil rigs (edge) and HQ (cloud). The edge has a UKG with equipment maintenance knowledge and can answer questions on-site quickly without internet. However, all edges report back usage data or new insights to the cloud UKG (which might consolidate learning). The cloud UKG sends updates like "new safety procedure" to all edges. This is orchestrated possibly by an update manager microservice.
- For the MVP, a hybrid model might not be fully built out, but the architecture is such that it is feasible. That means designing the system to separate configuration/data from code. The knowledge graph could be exported/imported easily. Perhaps using container images for code and separate volume or file for data makes it simple to ship updates.
- Another hybrid angle is **cloud-assisted edge computation:** If an edge device is resource-limited, it could optionally call the cloud for heavy computations (like generating a large embedding or running a complex simulation layer). For instance, edges could run up to a certain layer and ask the cloud to do intensive refinement, then return the result. This requires reliable connection though, so it's more like a thin-client approach which is less true offline.

In all cases, **security and access control** remain consistent. The deployment model influences where data resides (cloud vs local), but the system uses encryption (TLS) and authentication either way. In cloud, likely using OAuth tokens; in edge, maybe a simpler user/password or even no auth if it's a closed network.

From a DevOps perspective:

- Provide **deployment scripts/manifests** for each scenario. Helm charts or K8s manifests for cloud, Docker Compose for single-node edge, and documentation on how to connect edge to cloud if hybrid (like how to configure the cloud endpoint and API keys).
- Ensure the system is configurable via environment variables or config files so that, for example, one can turn off certain modules in edge or point the sync to a cloud URL, etc.
- Keep the codebase unified; these are just different configurations.

The MVP demonstration might not implement a full two-way sync, but would outline how an edge update could be done (perhaps via an export/import feature of the knowledge base). The question specifically lists these models to see that the architecture is flexible to handle them, which we have addressed.

In conclusion: - **Cloud-first:** Full-featured, scalable deployment with all data centralized. - **Edge/offline:** Self-contained instance with pre-loaded data for secure or remote operation. - **Hybrid:** A combination where edge nodes operate independently but receive updates from (and possibly send data to) a central cloud node, ensuring up-to-date knowledge everywhere without needing constant connectivity.

Security and API Design

Security is paramount in UKG, given it deals with potentially sensitive enterprise knowledge and is meant to provide authoritative answers. The system incorporates multiple layers of security: authentication, authorization, encryption, and auditing. Additionally, the API is designed with clear routes and roles, ensuring that each endpoint is access-controlled and that the system can be extended safely.

Authentication and Identity: UKG supports OAuth 2.0 authentication for user-facing clients and API key authentication for system-to-system integration. In a typical setup: - **User Login (OAuth2):** Users (e.g., employees of a company) authenticate via an identity provider. For MVP, this could be a simple username/password login managed by UKG's own database, or integration with an external IdP (like Azure AD, Okta) using OAuth2/OIDC. Upon login, the user receives a JWT (JSON Web Token) or bearer token to include in subsequent requests. The FastAPI backend has OAuth2 middleware that verifies these tokens on each request, ensuring the user is who they claim ¹⁴³. Multi-factor authentication (MFA) is recommended for added security ¹⁴³ (though might depend on external IdP if using one). Passwords (if any) are stored hashed (e.g., using bcrypt). - **API Keys and Service Auth:** For other systems (like if UKG API is called by an external process or as a tool), API keys can be issued. These are long random tokens that map to a service account with specific permissions. The system can require these keys to be passed in headers. All keys and secrets are stored securely (in environment vars or a vault, not in code).

Authorization (Access Control): Not all users should access all data or functions. UKG implements **role-based access control (RBAC)**. Users and API keys are associated with roles (or groups) that determine what they can do: - **Basic roles** might be: *Viewer* (can query and read data), *Contributor* (can add new knowledge or annotations), *Admin* (full rights including managing users). - **Row-level security:** In knowledge graph context, this could mean if some data is confidential, only certain roles or users can access those nodes. For example, there may be knowledge pillars that are restricted. The system can enforce this by checking user roles against node metadata (each node might have an access level tag). E.g., a node could be marked "Classified" and only users with a "Classified Access" role can retrieve it. If a query tries to retrieve something the user isn't allowed to see, the engine will omit that from results (or redacts it). Fine-grained control like this was alluded to: "*fine-grained access control – controlling who can see which nodes/relations*" ¹⁴⁴. - The API endpoints themselves require certain scopes/permissions. For instance, the `/admin/*` routes might require an admin token. FastAPI can use dependency-based security, where each route declares required roles, and a dependency function verifies the user token has that role.

Secure Communication (TLS): All network traffic is encrypted using TLS (HTTPS). In cloud deployments, this is a must – typically handled by the ingress or load balancer. For on-prem or closed network, TLS is still recommended especially if multiple components (like a separate DB server) communicate – e.g., ensure the database connection is over SSL. If using WebSockets (perhaps for streaming responses or a real-time chat

UI), those are done over `wss://` (which is essentially TLS for websockets). Certificates are managed and rotated according to best practices. Internally, if microservices talk, we can either rely on cluster network security or also use TLS service mesh for internal encryption.

Audit Logging: Security includes auditing sensitive actions. UKG should log authentication events (logins, failures), data changes (who added or modified a knowledge node), and administrative actions. These logs should be stored securely and monitored for anomalies (like an account accessing an unusually broad set of data). This ties into compliance if needed (the mention of FedRAMP, NIST standards in docs indicates awareness of auditing ¹⁴³). For MVP, logging to console or file is fine; in production, integrate with a SIEM (Security Info and Event Management) system.

Key API Routes: The API is organized with clear routes for major functionalities. Assuming a base path of `/api`, some key endpoints are:

- **Query Endpoint:** `POST /api/query` – The primary endpoint to query the knowledge graph/simulation. The client sends a JSON with the query text and perhaps optional parameters (like desired persona or format). The server authenticates the user, then processes the query via the Query Engine and Simulation Engine, and returns the result. The response would include the answer, and possibly structured data like confidence and citations. This endpoint would be accessible to any authenticated user with query privileges. It could also be a WebSocket or streaming endpoint if we implemented incremental answer streaming (not mandatory for MVP).

- **Knowledge Base Management:** e.g.

- `GET /api/nodes/{id}` – retrieve a knowledge node by its coordinate or ID (with details like content and metadata). Useful for drill-down or if front-end wants to show the full text of a node that was cited.
- `POST /api/nodes` – add a new knowledge node (with appropriate JSON body containing coordinate, content, metadata). Protected to only allow those with a curator or admin role.
- `PUT /api/nodes/{id}` – update a node's content or metadata. Also restricted.
- `GET /api/nodes?query=...` – possibly a search endpoint for nodes (this might use vector search or simple filters). For UI auto-complete, etc.

- **System Configuration:** `GET /api/config` – Admin-only endpoint to retrieve current system config (like confidence threshold, recursion limit, enabled modules, etc.). `PUT /api/config` – to update configuration settings at runtime (like toggling a persona or adjusting an algorithm parameter). Many config changes might require a restart, but some could be dynamic.

- **Persona Management:**

- `GET /api/personas` – List the configured personas (like "DomainExpert: axis8, description X; IndustryExpert: axis9, description Y; ..."). Useful for UI if it allows the user to direct a question to a specific persona or to show which personas are active.
- `POST /api/personas` – (if extension allowed) add a new persona definition. E.g., a custom persona that a developer creates for a new role. This would involve specifying which axis/role it maps to and perhaps some prompt or rules. Likely admin only.

- Possibly `PUT /api/personas/{id}/enable` or some toggle if one wants to temporarily disable one of the four personas (maybe not needed in MVP, but good for debugging).

- **User Management (if not external):** If UKG manages its own users:

- `POST /api/users` (admin creates a user),
- `GET /api/users` (list users, admin only),
- and `DELETE /api/users/{id}`, etc.
- And Auth endpoints: `POST /api/login` for obtaining token (if not using external OAuth flow).

- **Logs/Audit:** Possibly an endpoint like `GET /api/audit?since=...` for admin to fetch audit logs (or this might be outside the scope of API, could just rely on log files).

These routes are all documented in an OpenAPI spec generated by FastAPI, making integration easier for clients.

Security of API Routes: - The `query` endpoint is user-level and is heavily used; ensure it sanitizes input (though since we don't do raw SQL or shell commands with it, the main risk is prompting injection for the LLM if any – but here it's structured algorithms, not direct LLM, so okay). But still, we limit input length and validate it to avoid abuse (like extremely large queries causing huge memory use). - Rate limiting might be implemented on a per-user or per-key basis to prevent denial-of-service by flooding queries. - The management endpoints are admin-only and should additionally perhaps require strong auth (maybe only allow key-based auth from known IPs for config changes). - All state-changing requests (new node, delete node) should be CSRF-protected if used from web (with same-site cookies or tokens) and double-confirmed if critical (maybe not needed in API context if our UI just calls API).

Encryption and Data Security: Beyond communication encryption, data at rest encryption is considered. If using cloud managed DB, enable encryption at rest. If on server, use disk encryption especially for any sensitive data store (or at least ensure backups are encrypted). The knowledge graph might contain sensitive info; if it's in Neo4j or Postgres, ensure those files are secure. Also if vector DB has personal data embedded, treat that carefully.

Compliance: If targeting enterprise, align with standards: e.g., - Use **FedRAMP Moderate/High** controls if deployed for US gov (which implies a lot of security controls – but MVP just notes it). - Follow **NIST 800-53** controls, **SOC 2** principles for logging and change management ¹⁴³. - Privacy: if any personal data in knowledge graph, ensure GDPR compliance (right to be forgotten etc. – though UKG likely deals with company knowledge, not personal PII, unless configured so). - These aren't functional in MVP but the design acknowledges them (e.g., fine-grained access control for privacy ¹⁴⁴, audit logs, etc., showing we're building with compliance in mind).

Example API usage flow: A user logs in via the front-end. They get a JWT. The front-end calls `GET /api/personas` to display which expert perspectives are being used. The user types a query "What are the requirements for X?", the front-end calls `POST /api/query` with JSON `{"question": "What are requirements for X?"}` and the JWT in header. The backend authenticates, authorizes (say the user is allowed to query Pillar X data), runs the simulation, and returns `{"answer": "...", "confidence": 0.998, "sources": [{"id": node1, "title": "...", "coordinate": "FAR.1.2.3"}, ...]}`.

The front-end shows the answer and might list sources with clickable titles. If the user clicks a source, the front-end calls `GET /api/nodes/FAR.1.2.3` to get full text or details of that node and then displays a modal with that info. All these requests are checked for auth and logged in the audit trail.

Security Testing: We would test the API for common vulnerabilities: - Ensure no direct SQL queries from user input, to prevent injection (likely not, since using ORMs or safe drivers). - Test with fuzzed query input to ensure the system handles unusual inputs gracefully (no crashes or huge slowdowns exploitable). - Try permissions: e.g., normal user trying to access admin route yields 403. - If using JWT, ensure proper expiration and maybe use refresh tokens if long sessions needed. - Protect against replay (enforce token expiry, use HTTPS to avoid eavesdrop). - Possibly implement **Content Security Policy** headers on responses to mitigate XSS in web (though our front-end likely a separate static app). - Use secure cookies (HttpOnly, Secure, SameSite) if any session cookies are used.

Overall, the **API design** aims to be clear, RESTful, and secure by default. It provides the necessary endpoints for querying data (the main use case), as well as managing the knowledge base and system for authorized personnel. Every endpoint requires valid auth, all data transfers are encrypted, and the principle of least privilege is applied so users/apps only get access to what they need. By following these practices, the UKG MVP's API can be safely used in production environments.

Developer Guidelines and Extension Strategy

The UKG MVP is built to be **extensible and customizable**, recognizing that organizations may have unique knowledge domains or may want to enhance the system's reasoning capabilities. This section provides guidance for developers on how to extend the system – whether it's adding a new axis, integrating new knowledge content, introducing custom personas, or plugging in new algorithms.

Extending the 13-Axis Schema: The 13 axes defined are meant to be comprehensive, but there could be cases to adjust or expand them. Some possible extensions: - *Custom Axis Values*: More frequently, extension will mean adding new values or codes within an existing axis. For instance, adding a new Pillar domain or a new industry sector. The system is designed to accommodate that by updating the taxonomy data. Developers can add a new Pillar by assigning it the next number (or a specific number if integrating with an external taxonomy) and updating the relevant mapping (so that PillarName \leftrightarrow PillarID). Because axes are mostly data-driven, adding values doesn't require code changes – just insertion in the database or config files. The coordinate generator will then use those. - *Adding a Completely New Axis*: If truly needed (say an Axis 14 for a dimension not originally included), the architecture can theoretically support it (the coordinate vector concept can extend) ¹²⁰ ¹²¹. This would be a non-trivial change: developers would need to update data structures (perhaps making coordinate a 14-tuple), update any logic that iterates axes, and likely retrain/adapt any classification KAs that assumed 13 axes. It's doable but heavy. A more realistic scenario is repurposing an existing unused axis if any (some axes might not be fully utilized in certain use cases, those could be co-opted for another meaning). - *Perspective of Multi-Perspective Workflow*: The system's design in documentation even hints at "*additional axis layering*" in future for AGI modular fusion ¹²⁰, so they planned the possibility of more axes. If new axes are added, ensure: - The coordinate schema table is updated (like include new axis name, syntax rules). - The Query Engine knows how to classify for that axis (might need new ML model or rules). - The storage can handle it (if using separate fields, add a column, if using dynamic structure, just start using it). - Backward compatibility: older coordinates without that axis might need default or null for that axis.

Injecting New Knowledge or Domains: One of the most common extensions will be loading new knowledge into UKG. This could be: - **New Document Ingestion:** e.g., a company wants to load its internal policy documents or a new set of regulations. To do this: - **Coordinate Assignment:** Determine how the new content fits into the 13-axis scheme. For instance, if loading a company policy, maybe Pillar = Company Policies Pillar, Sector = the relevant dept, etc. Or if it doesn't fit existing Pillars, you might create a new Pillar for "Company Policy". - **Data Ingestion Tools:** Use or build an ingestion script. The developer can write a script that reads the document, breaks it into sections (hierarchy), and calls internal APIs or methods to create nodes and edges. We likely have utilities for this (like a method `create_node(coordinate, title, content, meta)` and `create_edge(source, target, type)`). Ensuring that for each section we also specify any cross-links (like link this policy section to an external regulation if applicable). - **Embeddings:** After adding text nodes, call the vector index service to generate and store embeddings. Possibly we provide a batch embedding function. - We'll want to update search indexes - if we have any full-text search separate from vector, update that too. - If ingestion is large, consider doing it offline or gradually, rather than via API one by one (performance). - The system does not need a full restart to include new knowledge, if ingestion goes through the API or a script using the system's libraries, the new nodes go into the graph in memory and into the DB. The Query engine and others inherently will consider them next query (especially if coordinates align and embeddings are in place).

UKG's framework may include an **import pipeline** with configuration (some docs mention knowledge ingestion pipelines with versioning ¹⁴⁵). So developers could configure an ingestion YAML that maps a data source to axes. For MVP, though, likely manual or scripted.

If developers want to maintain their knowledge base in an external source (like a database or CMS) and sync it to UKG, they could write an integration that listens for external changes and calls UKG's API to update or vice versa. The key is always providing the coordinate and enough context so UKG knows where it fits.

Custom Personas and Roles: Some organizations might want to simulate more than four personas, or different ones. While the core design is "Quad" Persona, the architecture can support instantiating additional personas if needed. For example, adding a "*Customer Perspective*" persona for queries dealing with customer experience (not exactly Axis 8-11 maybe, but one could map it to Axis 8 or introduce Axis 8 composite). - **To add a new persona:** - Decide which axis and value it corresponds to. If the persona is a variant of existing ones, you might just handle it within those (like treat it as a sub-role on Axis 8 or something). If it's something orthogonal, perhaps you use an axis that's free or double up on an axis meaning (not ideal). - Add a definition in the Persona configuration (like a persona id, name, axis mapping, any specific knowledge or biasing factors). Possibly implement a new Knowledge Algorithm if it needs special handling. - Integrate it into the Simulation Engine flow: likely, you'd modify Layer 3 to spawn 5 personas instead of 4. Also update aggregation logic to include it. - Make sure the Quad (now Quint?) persona coordination still works (the system might have assumed 4 in places, so those need generalizing). - Knowledge Algorithm KA-12 (Role Simulation) would need to be aware of the new persona's existence ¹⁴⁶ ¹⁴⁷. If our design is flexible, KA-12 might actually just loop through all configured personas, which would naturally include the new one if configured. - **Memory flow:** Additional personas would each have their own context copy. System resources scale accordingly. - If not adding an entirely new persona but **changing persona behavior** (like making the Industry persona think like a specific industry professional vs a generic one), developers can tweak the persona profile (maybe a prompt or a set of knowledge weighting). The design might have persona profiles declared in a config file that can be edited.

Plug-in Architecture for Simulation Algorithms: The simulation engine uses many Knowledge Algorithms (KAs) for its reasoning steps ⁹. We want developers to be able to add or override these algorithms, enabling custom logic or improvements. - The system could define an **interface for KAs** (in pseudocode, maybe an abstract base class with methods like `initialize(context)`, `execute(state)`, `produce_output()`). Each KA is assigned an ID (KA-1, KA-2, etc.) and the engine knows at which layer it's used. - To add a new algorithm, e.g., a "*Multilingual Reasoning KA*" to handle content in multiple languages (if the enterprise needs that): - Implement the class following the KA interface. Possibly specify which layer or sub-process it hooks into. - Register it in the system – could be via configuration (like list of active KAs and mapping to engine steps). The engine would then invoke it appropriately. - Alternatively, some algorithms might run in parallel to existing ones and provide additional input (like an ensemble). For example, if someone devises a new way to do consistency checking, they could add it to Layer 6. - **Validators and custom checks:** A common extension might be adding a domain-specific validator. Suppose an aerospace company wants every answer to go through an "ITAR compliance check" algorithm. They can develop a KA for that (maybe running in Layer 10 or as part of refinement) that scans the answer for any export-controlled info and flags it. They plug it in, and the refinement loop then includes that step. - The architecture could allow **plugins** in the form of Python modules that the system can dynamically load. For MVP, simpler is to have them integrated at build time, but providing a clear separation (e.g., a `plugins/` directory where additional KAs or persona definitions can be dropped in, and the system loads all found). - Another extension area: **tools integration** via MCP. If a developer wants UKG to use an external calculation engine for some queries, they could integrate that as an algorithm or as part of the chain. For instance, if a query requires heavy number crunching, a KA could call an external API or library to do that calculation and then return results into UKG's reasoning. Provided this is deterministic and doesn't break the in-memory paradigm too much, it's fine.

Modular Design: The system is meant to be modular. The UI has its component library, the backend has distinct modules (coordinate engine, simulation, data access, etc.), and even the knowledge content is modular by axes. Developers should adhere to these boundaries when extending: - For example, if adding a new axis, you mainly work in the coordinate engine module and data model, not entangle it with UI directly. - If adding a new UI feature (say a new dashboard or a new way to visualize reasoning traces), use the existing component library and patterns rather than making ad-hoc styles. - If customizing axes or personas, update documentation (possibly the design system or architecture docs) so future devs understand the new elements.

Extension Scenarios and How-To: - *Add a New Domain of Knowledge:* Define Pillar ID, insert Pillar node. Ingest documents under it. Possibly train or adjust the classification model for Query Engine to recognize queries related to that domain (or add keywords to rule-based classifier). - *Improve Answer Refinement:* Developer notices the writing style is too verbose. They want to implement a KA that concisely rephrases answers. They add that algorithm to the refinement steps. Or they integrate a language tool like Grammarly API as a check in Step 12. - *Scaling Up:* If moving from MVP to large-scale, a developer might want to swap out components (e.g., use a distributed vector search instead of local, or an LLM for generating explanations). Having a plugin mentality helps: e.g., we can plug in an LLM as KA-15 "Language Polisher" that takes the structured answer and rewrites it in nice prose. - *User Interface Customization:* This might involve adding new pages (maybe an Admin panel to manage knowledge nodes). Follow the design system: use the provided styles for tables, forms. If none exists for a needed pattern, extend the component library for that pattern (and perhaps contribute back for consistency).

Maintaining Consistency: When extending, it's important not to break the core logic. For example, if customizing an axis, ensure that you don't violate the coordinate syntax rules (because other parts parse them). Always test after extension: run queries that involve the new elements to make sure they integrate smoothly.

Developer Tools: Provide some tools to help extension: - Possibly a CLI tool that can validate the knowledge graph (to ensure all coordinates are valid, all references resolve, etc.). - Scripts for one-time tasks (like a script to generate a new persona scaffold or a script to import a batch of nodes from CSV). - Testing harness: encourage writing unit tests for new KAs or modules. The structure might allow simulating a single layer or persona in isolation for testing logic.

Community and Modularization: If external developers or open-source community were involved, a plugin system with clearly defined extension points fosters collaboration. E.g., someone could write a plugin to integrate a quantum computing module for one of the reasoning layers and share it.

Documentation: When extending, update the system documentation (like this document) accordingly. If a new axis is added, it should be documented for future devs and users. If new personas are added, update the System Overview or UI to reflect that (the UI might need to show more personas debating, etc.).

In conclusion, the UKG MVP is built not as a closed black box but as an **evolving framework**. It encourages developers to tailor it to their needs by: - Extending the knowledge model (axes and data), - Enhancing the reasoning pipeline (new algorithms, personas), - Integrating with other systems (via MCP or custom KAs), - Customizing the user experience (new visualizations or workflows), all while using structured interfaces and configurations that keep the core stable. With these guidelines, engineering teams can implement the system from scratch and confidently adapt it as their requirements grow, ensuring UKG remains a **universal and flexible knowledge framework** for their organization.

Sources: The design and features described above are based on the UKG technical documentation and white papers, which detail the 13-axis model [3](#) [61](#), the multi-layer simulation approach [5](#) [90](#), the Quad Persona engine [101](#) [102](#), and the integration points like MCP [107](#) [148](#). The UI/UX principles follow the Enhanced Design System guidelines [23](#) [28](#). Security and compliance considerations align with enterprise standards mentioned in the documentation [144](#) [143](#). This concludes the comprehensive overview and developer guide for the UKG MVP.

[1](#) [8](#) [79](#) [80](#) [144](#) Universal Simulated Knowledge Graph (UKG) – A Comprehensive White Paper.pdf
file:///file-Xzcfcf2z4gpGr713LhAgYY

[2](#) [3](#) [4](#) [33](#) [34](#) [35](#) [36](#) [37](#) [38](#) [39](#) [40](#) [41](#) [42](#) [43](#) [44](#) [45](#) [46](#) [47](#) [48](#) [49](#) [50](#) [51](#) [52](#) [53](#) [54](#) [55](#) [56](#) [57](#) [58](#) [59](#)
[60](#) [61](#) [62](#) [63](#) [64](#) [65](#) [66](#) [67](#) [68](#) [114](#) [115](#) [116](#) [117](#) [118](#) [119](#) [127](#) [128](#) Unified Coordinate System for UKG_USKD –
Technical Documentation.pdf
file:///file-1ryYgxyFf1XhtZ4CsBY5iE

[5](#) [6](#) [90](#) [94](#) [95](#) [96](#) [97](#) [98](#) [146](#) [147](#) Universal Knowledge Algorithm (KA) System – Technical Report.pdf
file:///file-XRkEpP4exQhaXsqo696fhc

7 9 69 70 71 72 73 74 75 76 77 78 81 82 83 84 85 91 92 93 125 Universal Knowledge Graph

Simulation Stack – 10-Layer Architecture Report.pdf

file://file-Ut87GALyFv3ccojTzDc632

10 11 32 86 87 88 89 99 100 101 102 103 104 105 106 107 108 109 110 111 112 113 120 121 122 123 124 145

148 Integrating the Model Context Protocol (MCP) into the Universal Knowledge Graph (UKG).pdf

file://file-NctsBAWUjZHTxVPpMGWwRr

12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 129 130 131 132 133 134 135 136 137

138 139 140 141 142 Universal_Knowledge_Framework_-_Enhanced_Design_System.pdf

file://file-2jENeHJfXvvjPjs34jL3co

126 143 Universal_Knowledge_Database_Mathematical_Framework_White_Paper.pdf

file://file-BFxXpagcJSTcz4TkSKp5u5